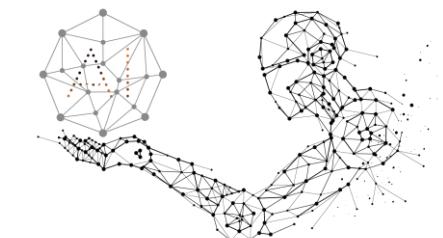


Applied Machine Learning

Lecture 6- A little Probability

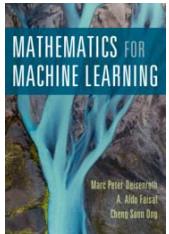


Hossein Homaei
Department of Electrical & Computer Engineering

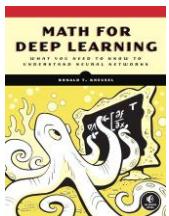


Some resources

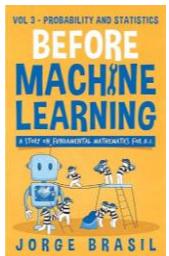
- Books



M. P. Deisenroth, A. A. Faisal, and C. S. Ong, *Mathematics for machine learning*. Cambridge University Press, 2020.



R. T. Kneusel, *Math for Deep Learning: What You Need to Know to Understand Neural Networks*. No Starch Press, 2022.



J. Brasil, *Before Machine Learning*, vol. 3, Probability And Statistics. 2024.



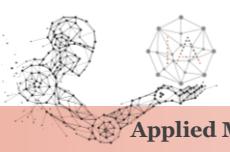
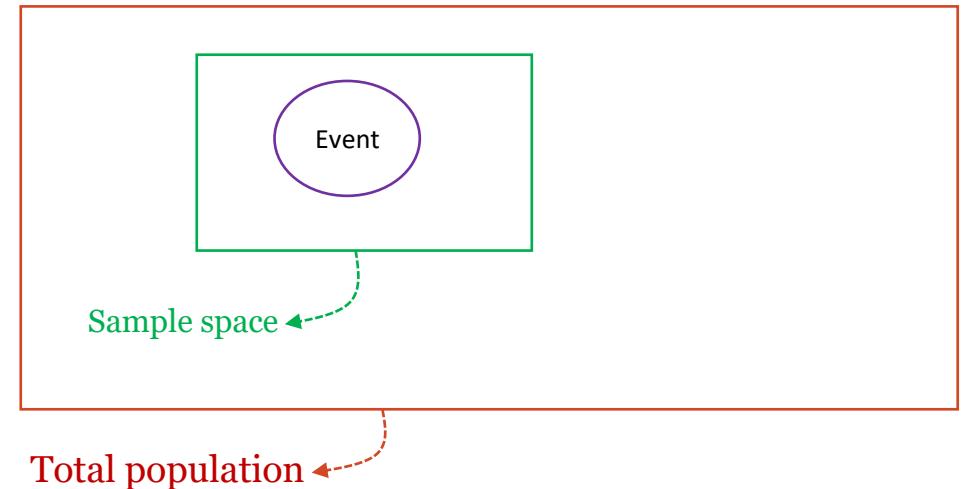
... Some resources

- Online
 - Probability & Statistics for Machine Learning and Data Science
 - Instructor: Luis Serrano
 - DeepLearning.AI
 - **A significant portion of this lecture's content is sourced from this course.**
 - CS229: Machine Learning- The Summer Edition 2019
 - Instructor: Anand Avati
 - Stanford University



Introduction

- Probability
 - A measure of how likely an event is to occur
 - $p = \frac{\text{size of the event}}{\text{size of sample space}}$
- Example
 - All children = population
 - Children of a school = sample space = 100
 - Event = children play soccer = 40
 - $p = \frac{40}{100} = 0.4$



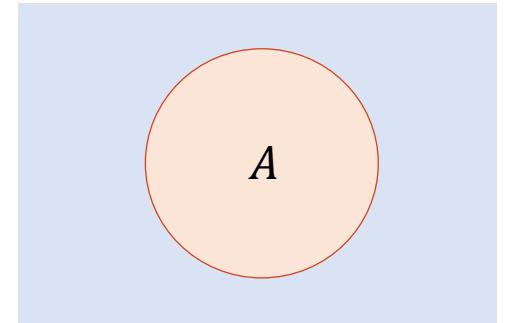
... Introduction

- Experiment
 - Any process that produces an outcome that is uncertain
 - Example: throwing a coin
 - $p = \frac{\text{size of the event}}{\text{total number of outcomes}}$
 - Examples
 - Roll a fair dice: The probability of obtaining 6= $p(6) = \frac{1}{6}$
 - Flip a fair coin twice: The probability of landing on heads twice= $p(HH) = \frac{1}{4} = 0.25$



Probability Rules

- The complement rule
 - $p(\bar{A}) = 1 - P(A)$
 - Roll a dice. Probability of not obtaining 6?
 - $P(\text{not } 6) = 1 - p(6) = 1 - \frac{1}{6} = \frac{5}{6}$



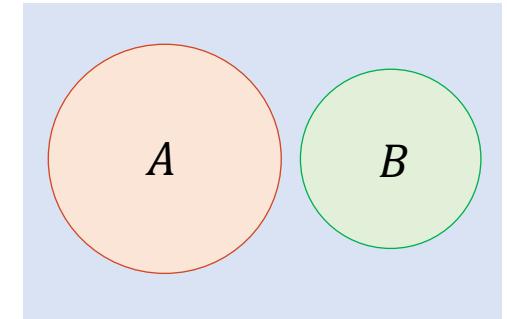
Venn Diagram



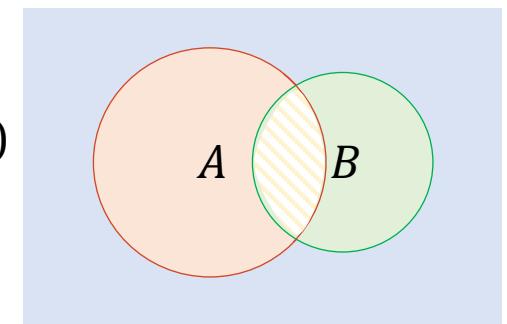
... Probability Rules

- Sum of probabilities
 - For **disjoint** events A and B , $P(A \cup B) = P(A) + P(B)$
 - Roll a dice. Probability of obtaining an even number or 5?

$$P(\text{even OR } 5) = P(\text{even} \cup 5) = \frac{3}{6} + \frac{1}{6} = \frac{2}{3}$$



- For **joint** events A and B , $P(A \cup B) = P(A) + P(B) - P(A \cap B)$
 - Roll 2 dices. Probability of obtaining a sum of 7 or a difference of 1?
 - Sum=7: (1, 6), (2, 5), (3, 4), (4, 3), (5, 2), (6, 1)
 - Diff=1: (1, 2), (2, 1), (2, 3), (3, 2), (3, 4), (4, 3), (4, 5), (5, 4), (5, 6), (6, 5)
 - Sum=7 and diff=1: (3, 4), (4, 3)
 - $P(\text{sum7} \cup \text{diff1}) = \frac{6}{36} + \frac{10}{36} - \frac{2}{36} = \frac{14}{36} = \frac{7}{18}$



... Probability Rules

- Product rule

- For **independent** events A and B , $P(A \cap B) = P(A) \times P(B)$
 - Flip a fair coin 5 times. Probability of landing in heads 5 times?

$$P(5 \text{ heads}) = \frac{1}{2} \times \frac{1}{2} \times \frac{1}{2} \times \frac{1}{2} \times \frac{1}{2} = \frac{1}{32}$$

The probability of B given A

- Conditional probability: For **dependent** events A and B , $P(A \cap B) = P(A) \times P(B|A)$

- Roll a dice twice. Probability that the first is 6 and the sum equals 10?

$$P(\text{first } 6) = \frac{6}{36} = \frac{1}{6}$$

$$P(\text{sum} = 10 | \text{first } 6) = \frac{1}{6}$$

$$P(\text{first } 6 \cap \text{sum} = 10) = \frac{1}{6} \times \frac{1}{6} = \frac{1}{36}$$

1,1	1,2	1,3	1,4	1,5	1,6
2,1	2,2	2,3	2,4	2,5	2,6
3,1	3,2	3,3	3,4	3,5	3,6
4,1	4,2	4,3	4,4	4,5	4,6
5,1	5,2	5,3	5,4	5,5	5,6
6,1	6,2	6,3	6,4	6,5	6,6



Discrete Distributions



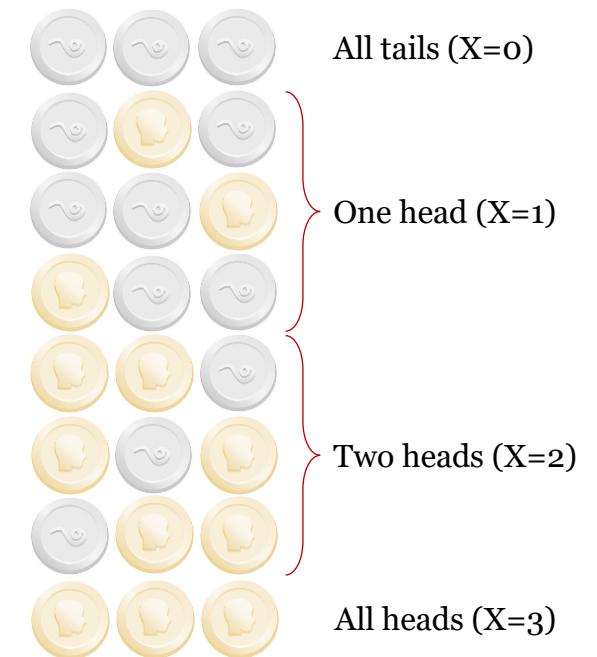
Random Variables

- A quantity that depends on random events (uncertain outcomes)
- Types
 - Discrete random variables: can take only a countable number of values
 - $X =$ number of heads in 10 coin tosses
 - $X =$ number of 1's in 7 dice rolling
 - Continuous random variables: can take values on an interval
 - Wait time until the next bus arrives
 - Quantity of rain (mm)



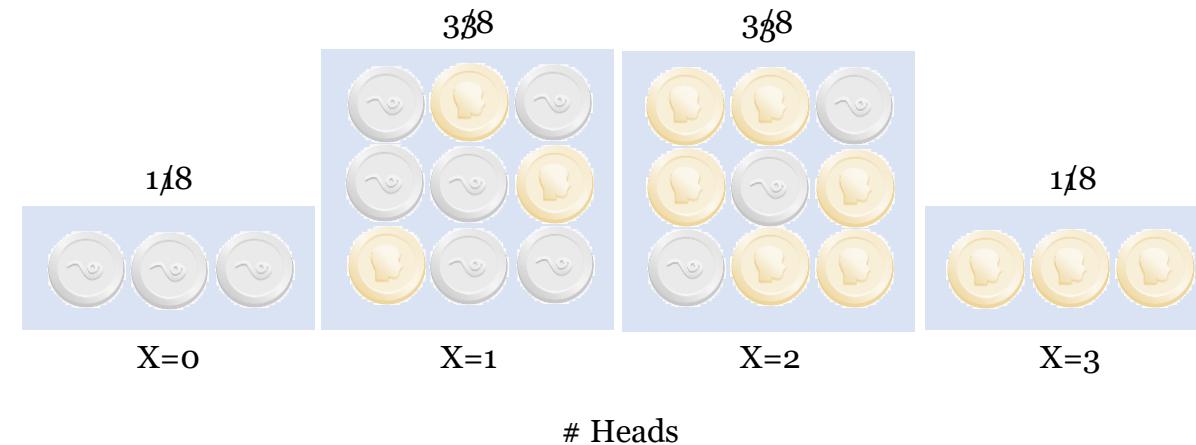
Probability Distribution

- Put all possible scenarios that can happen in the horizontal axis, and for each of them, specify the probability that it happens.
 - Example
 - Flip a fair coin 3 times
 - Random variable (X) = number of heads



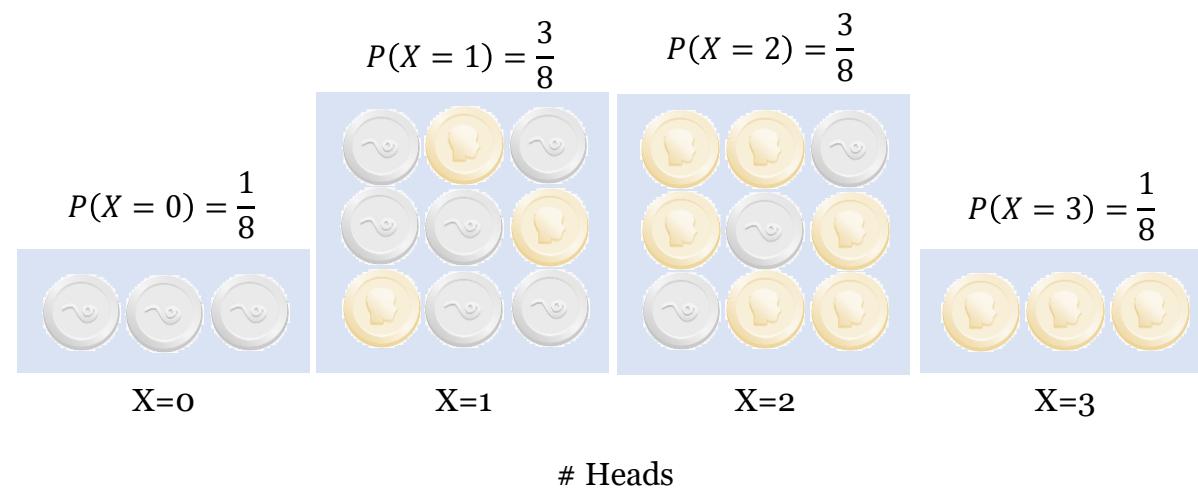
Probability Distribution

- Put all possible scenarios that can happen in the horizontal axis, and for each of them, specify the probability that it happens.
 - Example
 - Flip a fair coin 3 times
 - Random variable (X) = number of heads



... Probability Distribution

- Histogram interpretation
 - For each x from 0 to 3 (all possible values for random variable X), the probability that X is x



Probability Mass Function (PMF)

- What is it?
 - The function that represents how the probability distributes among all the possible values of the discrete random variable
- Notes
 - Use lowercase p to represent it.
 - $p_X(x) = P(X = x)$
 - All discrete random variables can be modeled by their PMF
- Properties
 - $p_X(x) \geq 0$
 - $\sum_x p_X(x) = 1$



Binomial Distribution

- Example: Flip 5 coins
 - What is the probability that 2 of them land in head?
 - For each flip, there is a $1/2$ chance to get heads or tails
 - For the fair coins, each of the outcomes occurs with the same probability $1/32$
 - There are 10 possibilities to get two heads out of 5.
 - Is there any general way to find the number of all possible combinations
 - Yes

$$10 = \frac{\text{Number of ways you can order 5 coins}}{\text{Number of ways you can swap 2 heads} \times \text{Number of ways you can swap 3 tails}}$$
$$= \frac{5!}{2! (5 - 2)!} = \binom{5}{2}$$

→ Binomial coefficient
Number of ways you can get 2 heads in 5 coin tosses



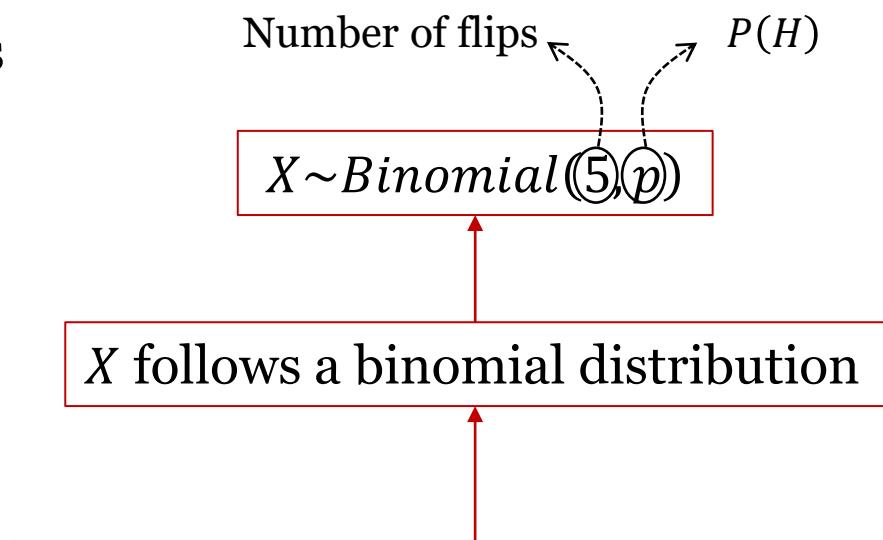
... Binomial Distribution

- Binomial coefficient
 - $\binom{n}{k}$
 - read as "n choose k "
 - Counts all the combinations in which you can have k heads in n coin tosses
- Property
 - $\binom{n}{k} = \binom{n}{n-k}$
 - This is the reason that the PMF of tossing a fair coin has a symmetrical shape



... Binomial Distribution

- Example
 - Looking for a general way to write the PMF for the number of heads in 5 coin tosses
 - X = Number of heads in 5 coin tosses
 - Suppose that the probability of heads is p
 - $P(H) = p$
 - Event $X = x$ means that x heads in 5 coin tosses
 - x can be 0, 1, 2, 3, 4, or 5
 - What is the probability of this event?
 - For one particular order:
 - Probability of seeing x heads = p^x
 - Probability of seeing $5 - x$ tails = $(1 - p)^{5-x}$
 - PMF → For all possible orders:
 - $p_X(x) = \binom{5}{x} p^x (1 - p)^{5-x}$ for $x = 0, 1, 2, 3, 4, 5$

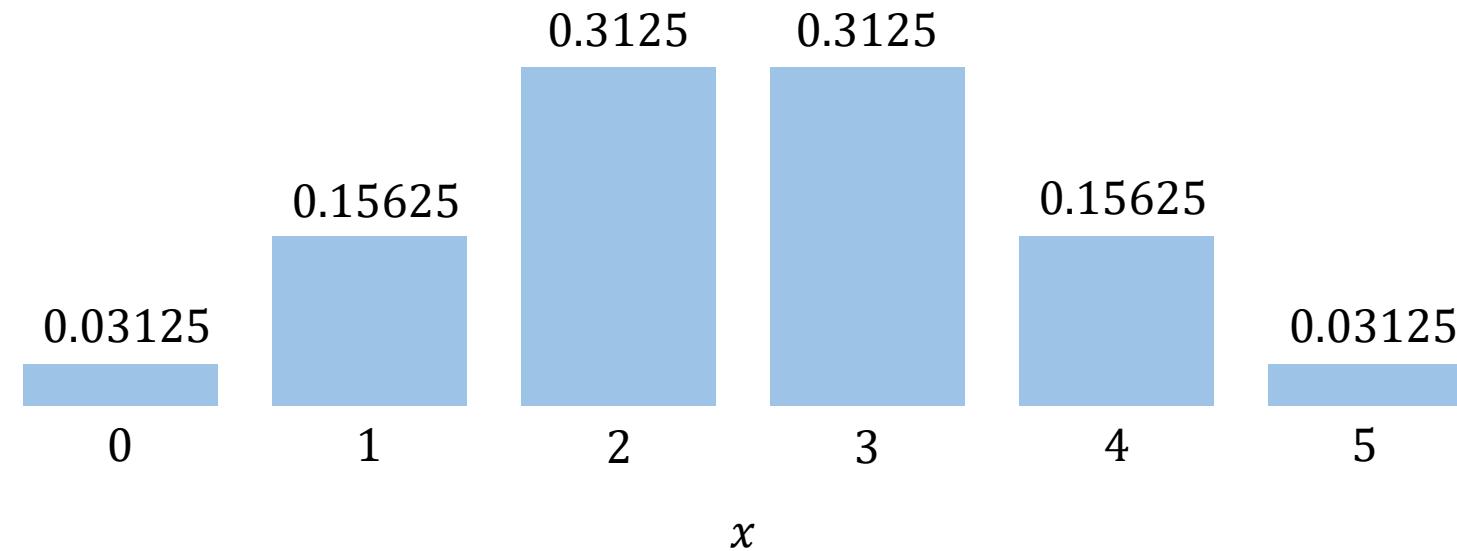


... Binomial Distribution

- Example- Symmetric

- $X \sim Binomial\left(5, \frac{1}{2}\right)$

- $p_X(x) = P(X = x) = \binom{5}{x} 0.5^x 0.5^{5-x}$

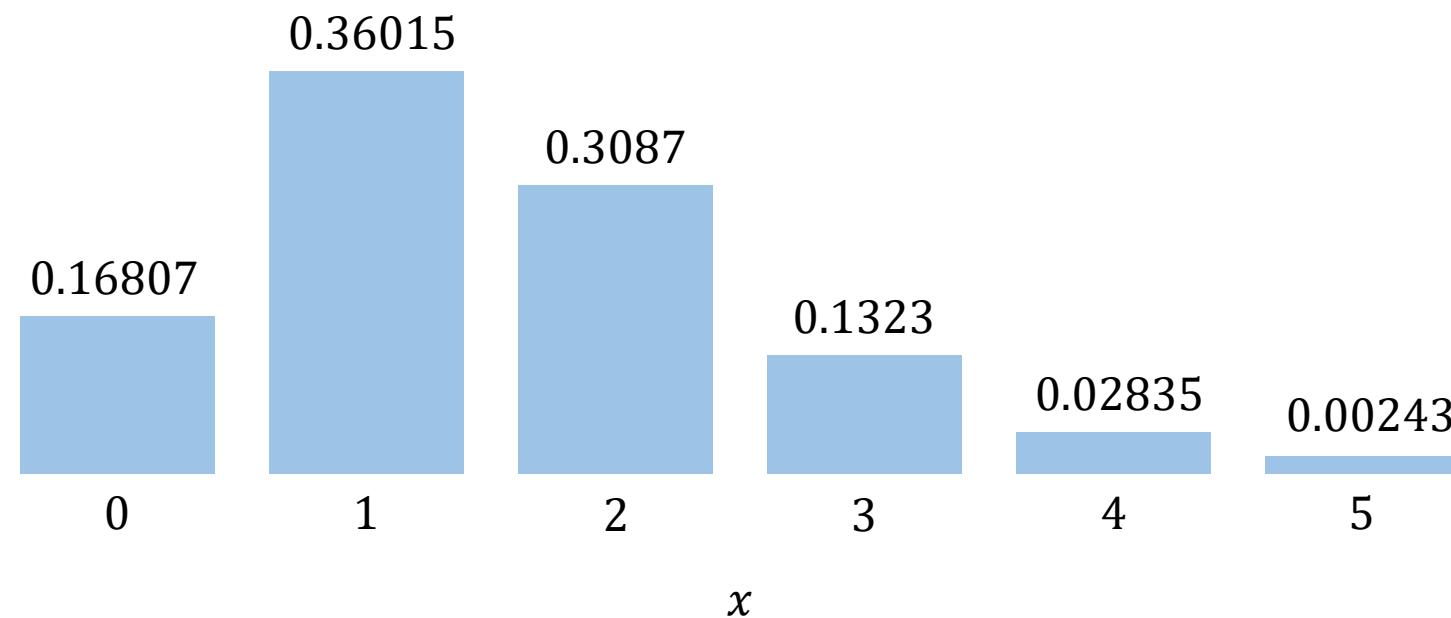


... Binomial Distribution

- Example- Asymmetric

- $X \sim \text{Binomial}(5, 0.3)$

- $p_X(x) = P(X = x) = \binom{5}{x} 0.3^x 0.7^{5-x}$



... Binomial Distribution

- General PMF formula
 - $X \sim \text{Binomial}(n, p)$
 - $p_X(x) = P(X = x) = \binom{n}{x} p^x (1 - p)^{n-x}$ for $x = 0, 1, 2, \dots, n$
 - n and p are called the parameters of the binomial distribution
- Number of heads in n coin flips
- $P(H) = p$
- Event $X = x$: x heads in n tosses



... Binomial Distribution

- Example
 - Rolling a dice 5 times
 - What is the PMF of getting ones?
 - We can assume that the dice is a biased coin
 - Get heads (one) with probability $1/6$
 - Get tail (others) with probability $5/6$
 - $X \sim \text{Binomial} \left(5, \frac{1}{6} \right)$

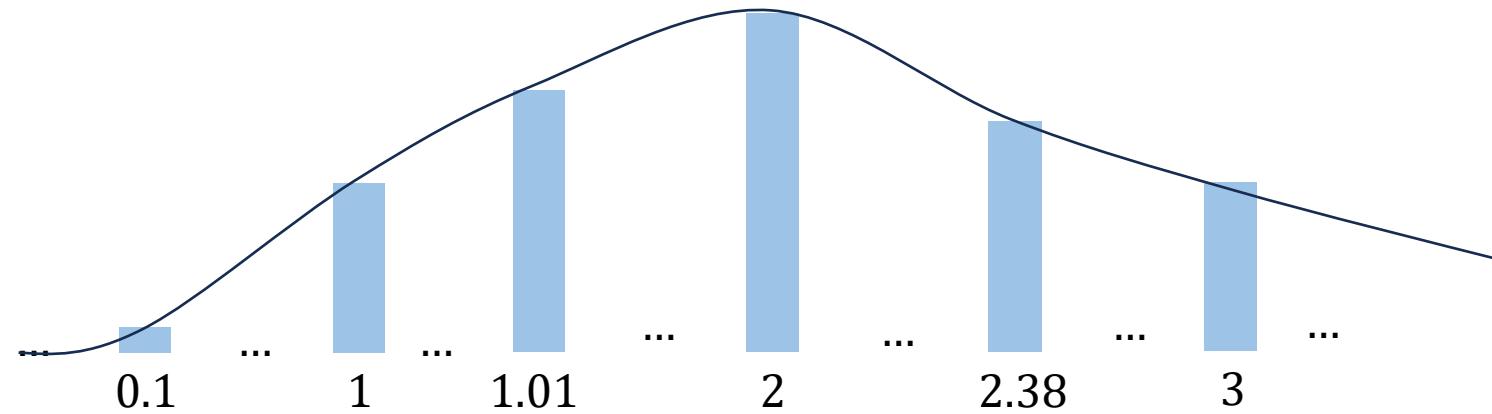


Continuous Distributions



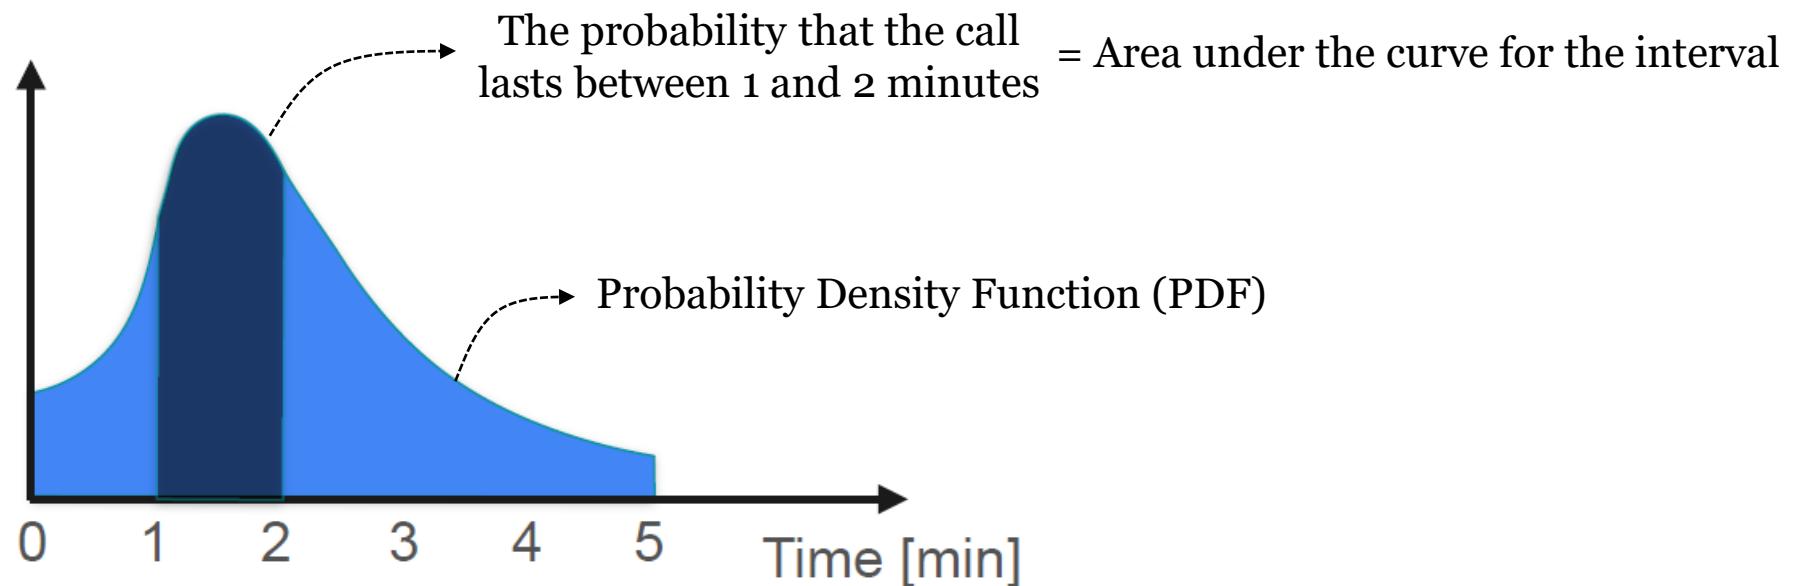
Introduction

- In the discrete distributions, the events always form a list (countable)
- In continuous distributions, the events are an interval
 - Example: waiting time on the phone call
- We cannot draw the probability distribution in the same way we did in PMF
 - Instead, we can draw it for intervals, e.g. 0.01 second intervals
 - Make the intervals more granular to get tiny intervals
 - For infinitely many skinny bars we get a curve



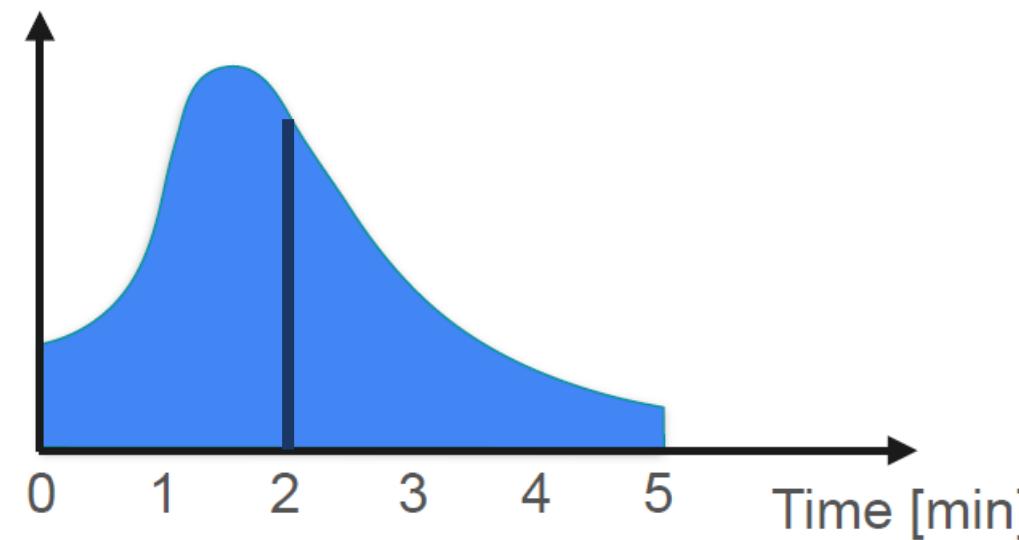
... Introduction

- Discrete variable
 - Sum of heights equals 1
- Continuous variable
 - Area under the curve equals 1



... Introduction

- What is the probability that the call is exactly 2 minutes?
 - Zero!
- In continuous distributions, we can only think of probabilities between window interval, not at a particular single point



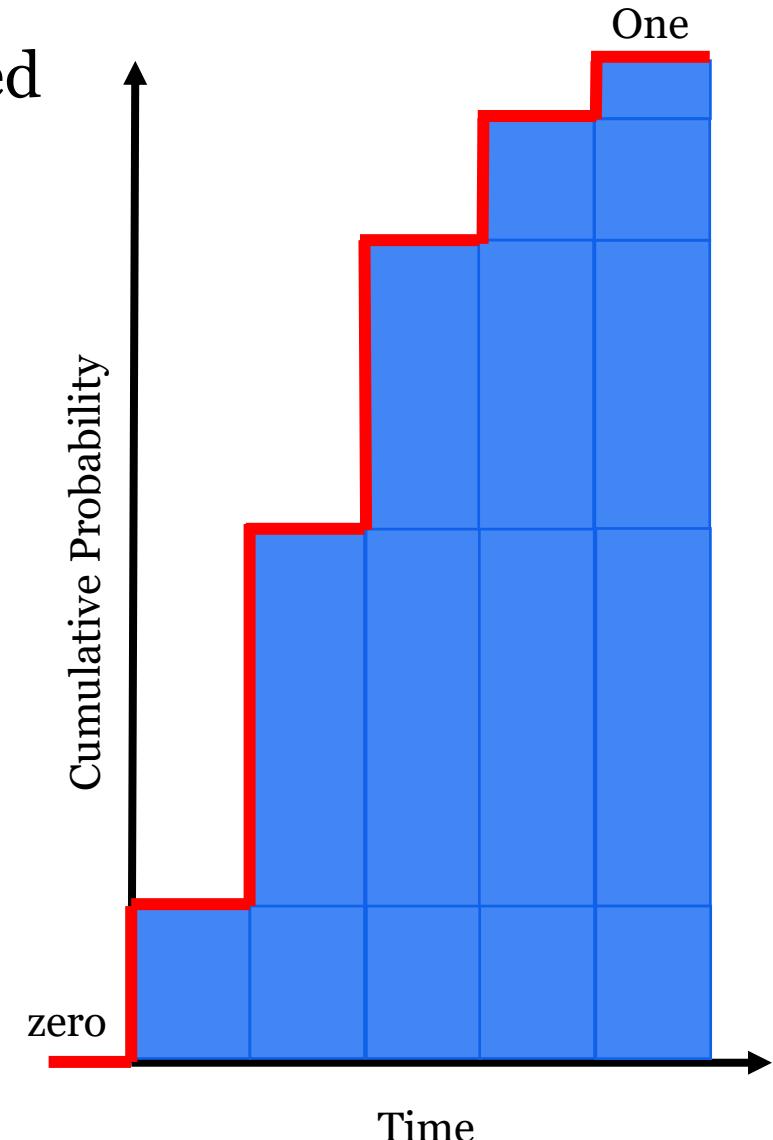
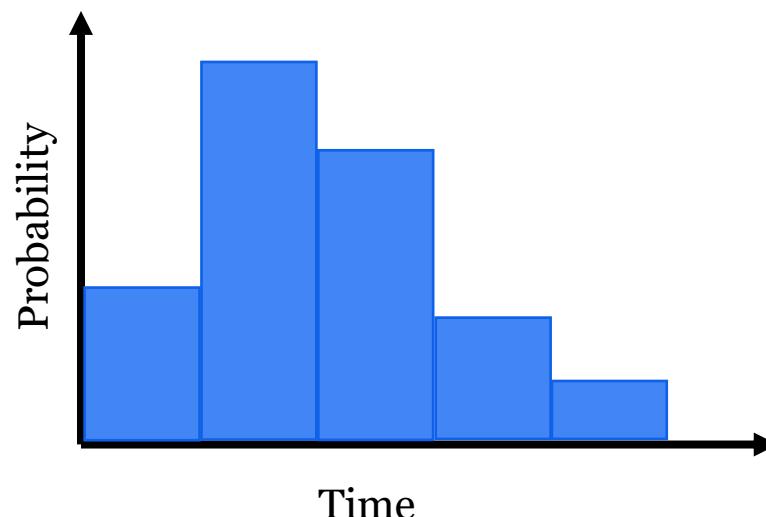
Probability Density Function (PDF)

- What is it?
 - The function that tells you the rate you accumulate probability around each point for continuous random variable
- Notation
 - Use lowercase f to represent it.
 - $f_X(x)$
 - $P(a < x < b) = \text{area under } f_X(x)$
- Properties
 - $f_X(x)$ is defined for all numbers
 - $f_X(x) \geq 0$
 - Area under $f_X(x)$ equals 1 ($\int f_X(x) = 1$)



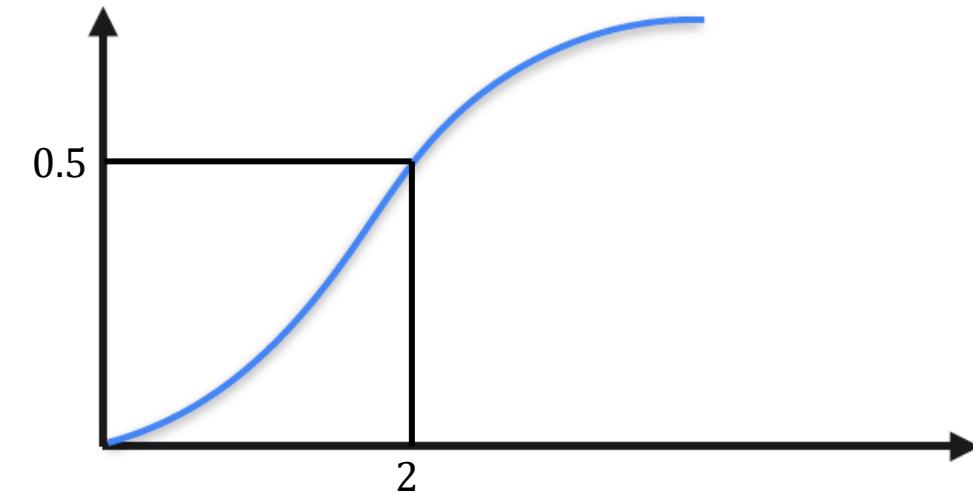
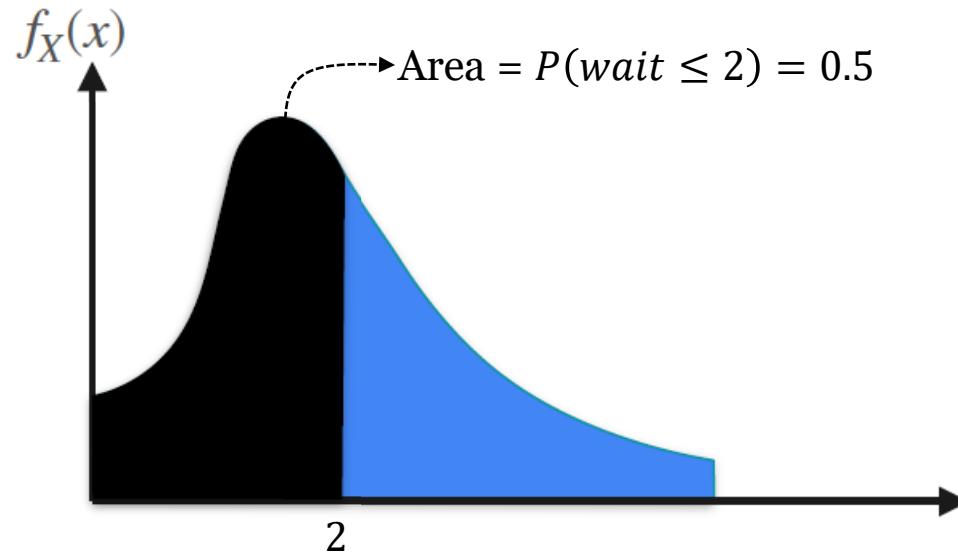
Cumulative Distribution Function (CDF)

- Specify what is the probability that an event happened before some reference point.
 - Example- call center waiting time
 - The discrete simplification
 - Property
 - Always start at zero and end at one



... Cumulative Distribution Function (CDF)

- Continuous distributions
 - Instead of adding the heights of bars, plotting the sum of the area under the curve



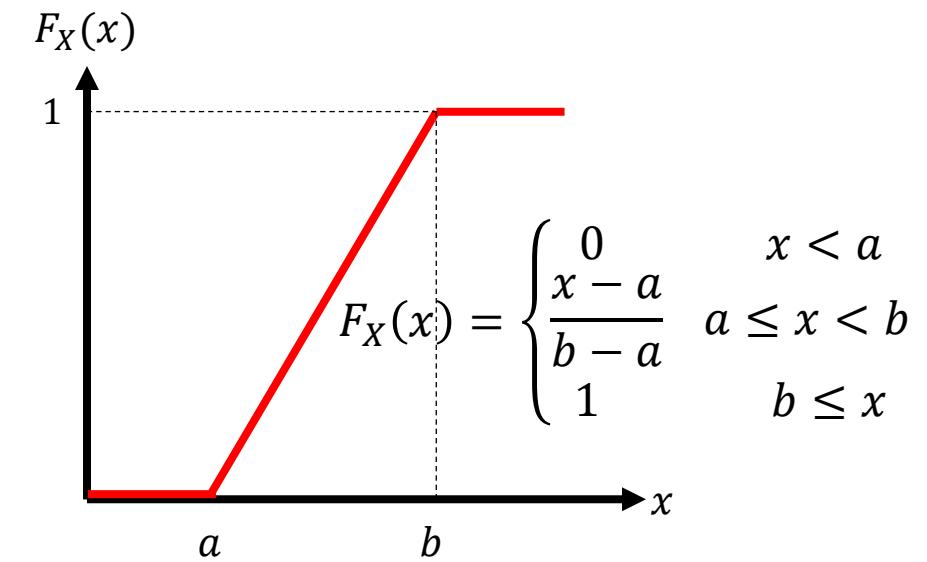
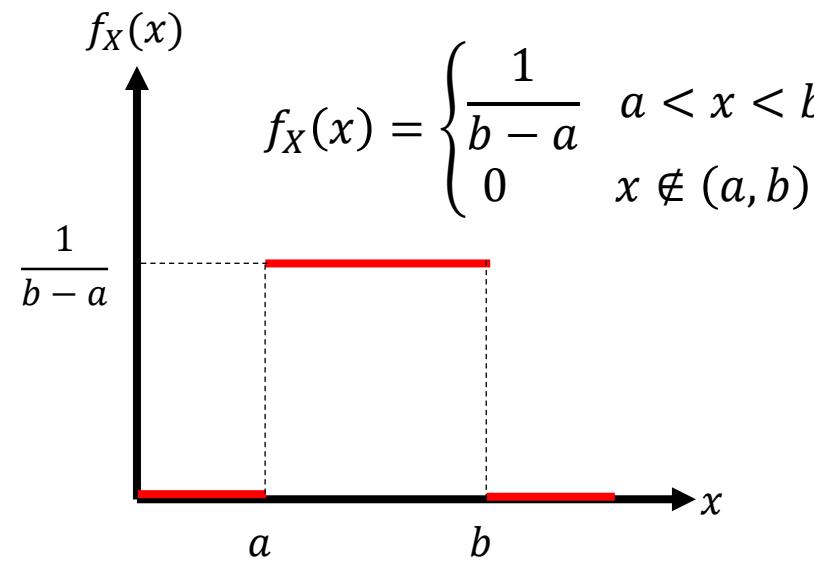
... Cumulative Distribution Function (CDF)

- More formal
 - The CDF shows how much probability the variable has accumulated until a certain value
 - Use uppercase F to represent it
 - $F_X(x) = P(X \leq x)$
 - Properties
 - $0 \leq F_X(x) \leq 1$
 - Left endpoint is 0
 - Right endpoint is 1
 - Can be at infinity
 - Never decreases



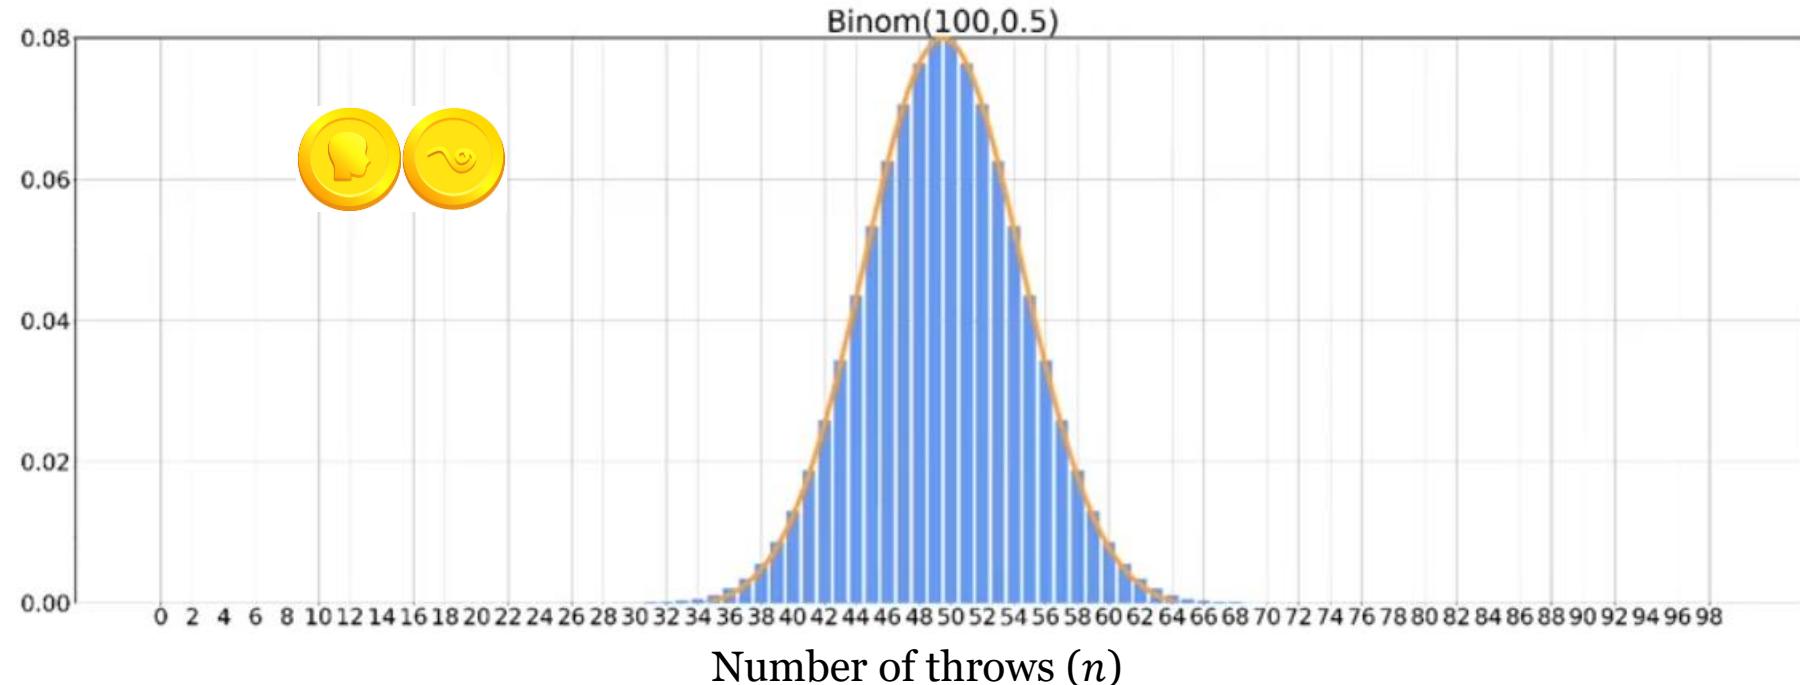
Uniform Distribution

- A continuous random variable can be modeled with a uniform distribution if all possible values lie in an interval and have the **same frequency** of occurrence
- Parameters
 - Beginning of the interval
 - End of the interval
- Formula



Normal Distribution

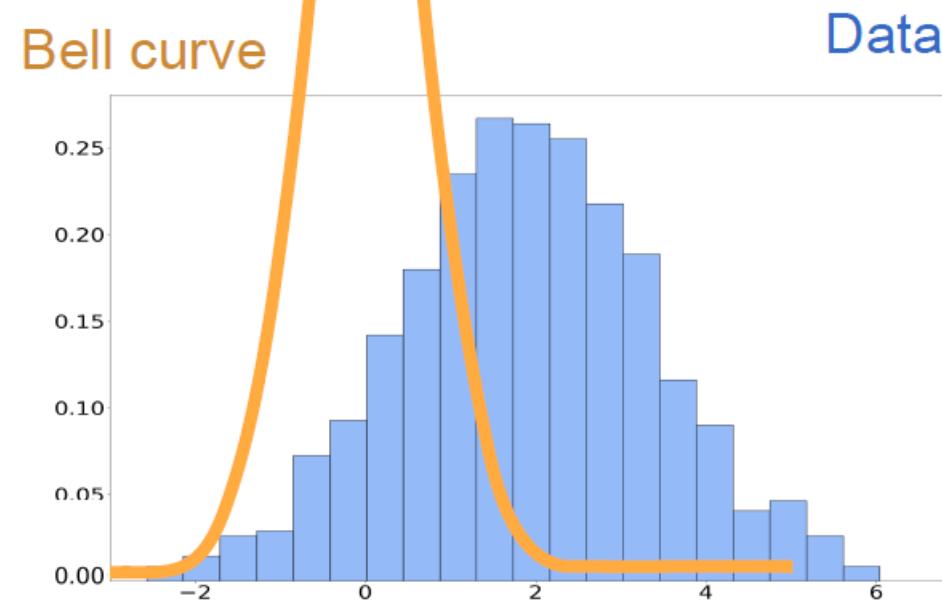
- Bell curve shape
- Also named Gaussian distribution
- The binomial distribution can be approximated by the normal distribution
 - Example- Throw a fair coin for n times, when n is very large



... Normal Distribution

- Formula- example

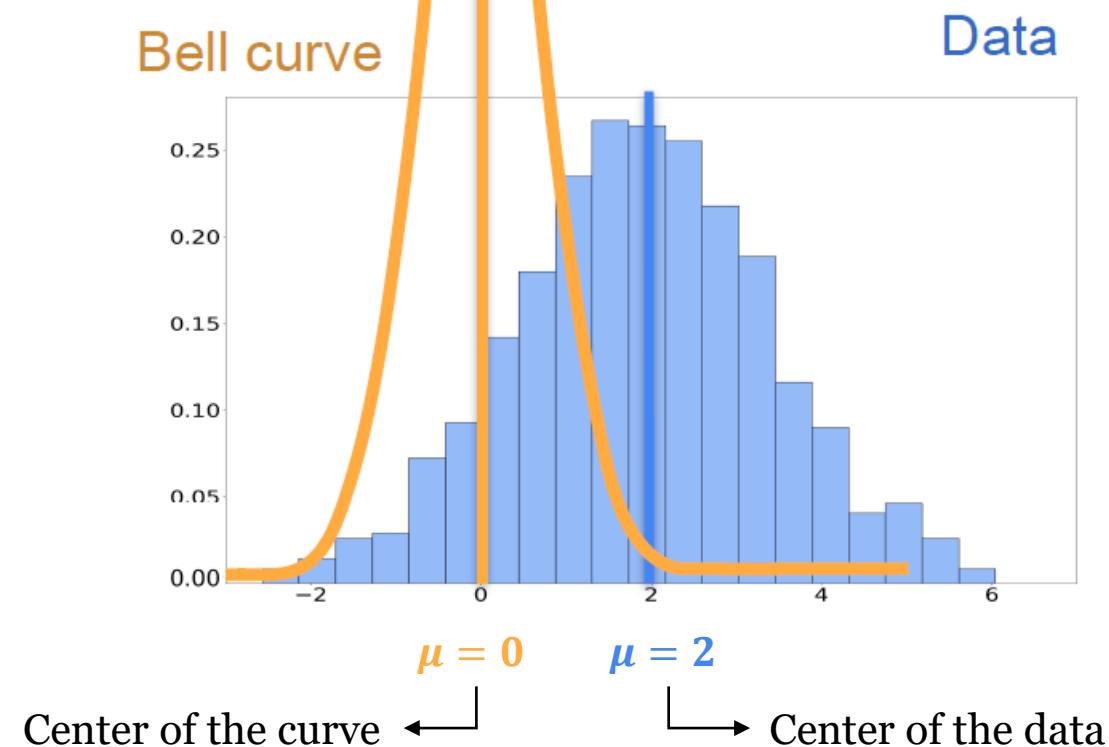
$$e^{-\frac{1}{2}x^2}$$



... Normal Distribution

- ... Formula- example
 - Not centered

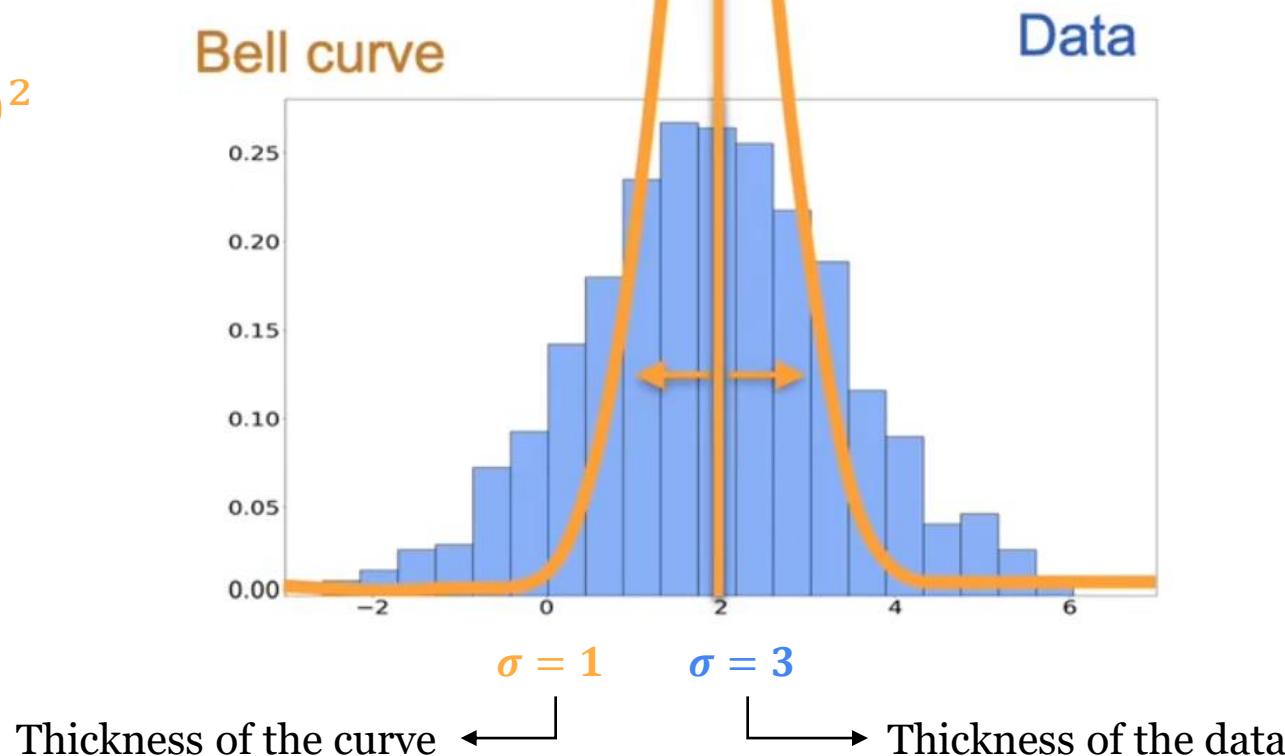
$$e^{-\frac{1}{2}x^2}$$



... Normal Distribution

- ... Formula- example
 - Skinny

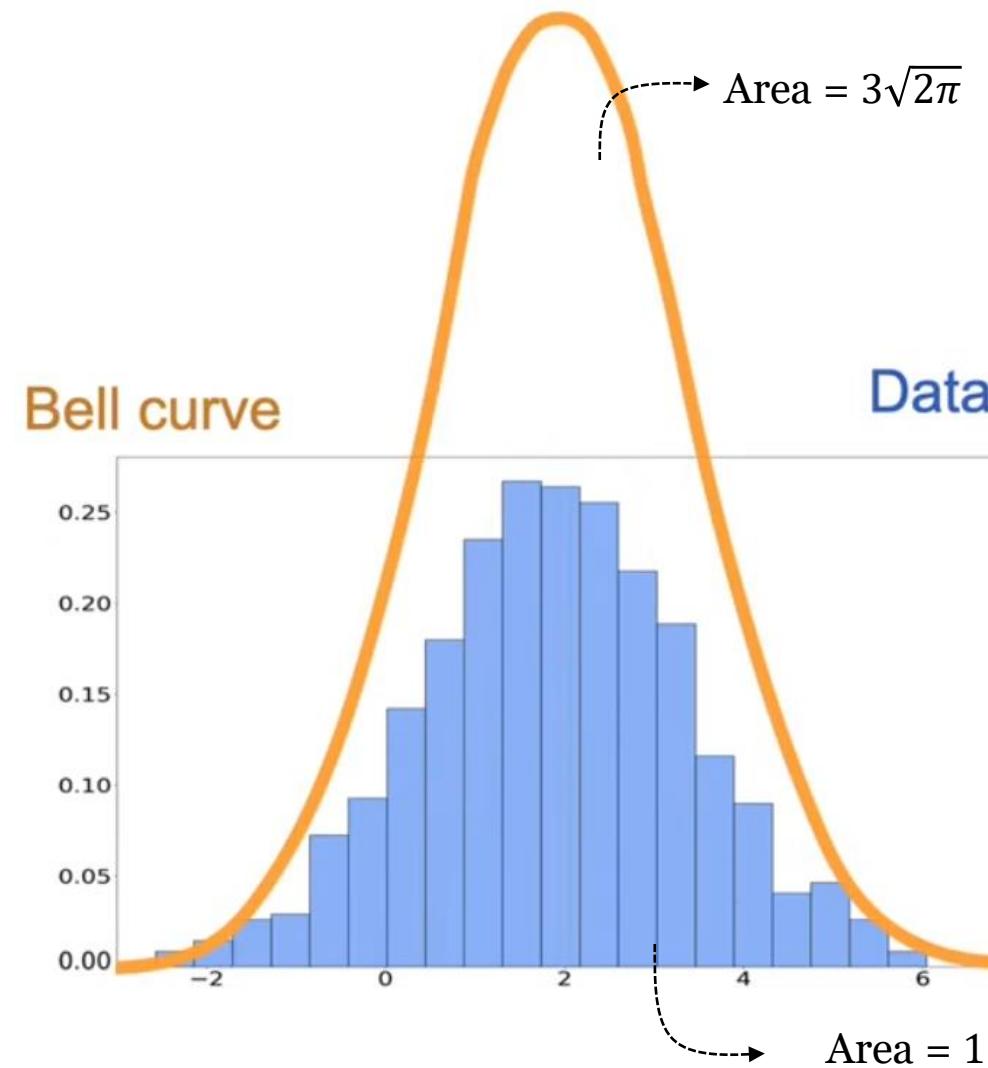
$$e^{-\frac{1}{2}(x-2)^2}$$



... Normal Distribution

- ... Formula- example
 - Different height

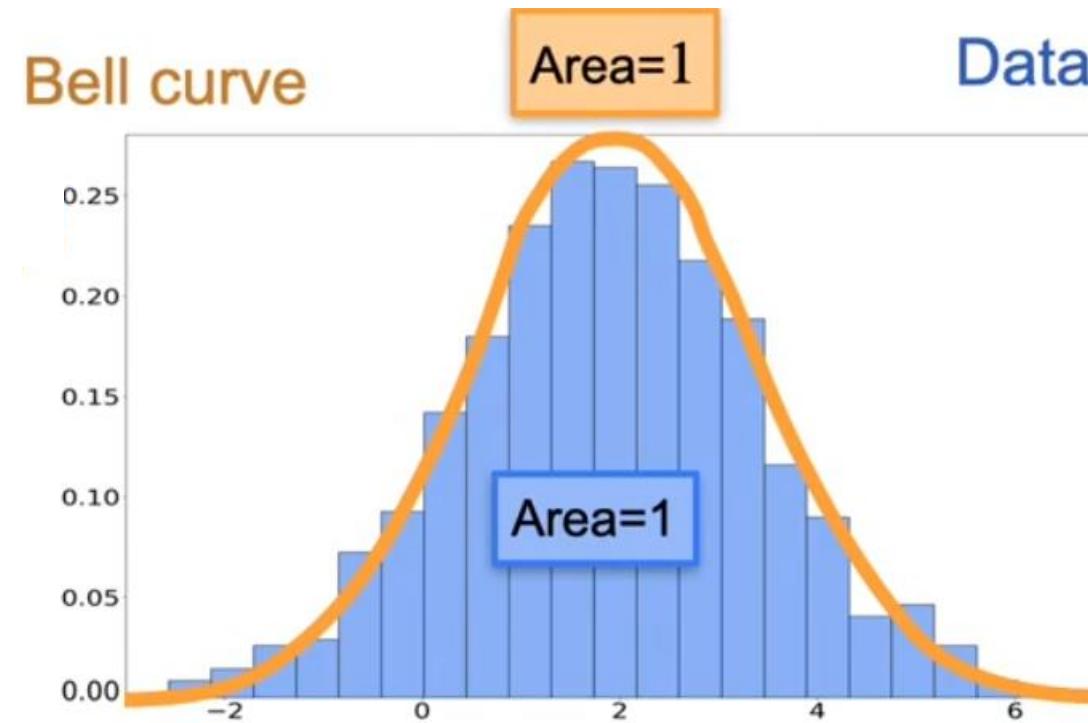
$$e^{-\frac{1}{2}\left(\frac{x-2}{3}\right)^2}$$



... Normal Distribution

- ... Formula- example

$$\frac{1}{3\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-2}{3}\right)^2}$$



... Normal Distribution

- Formula- general

$$\frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$

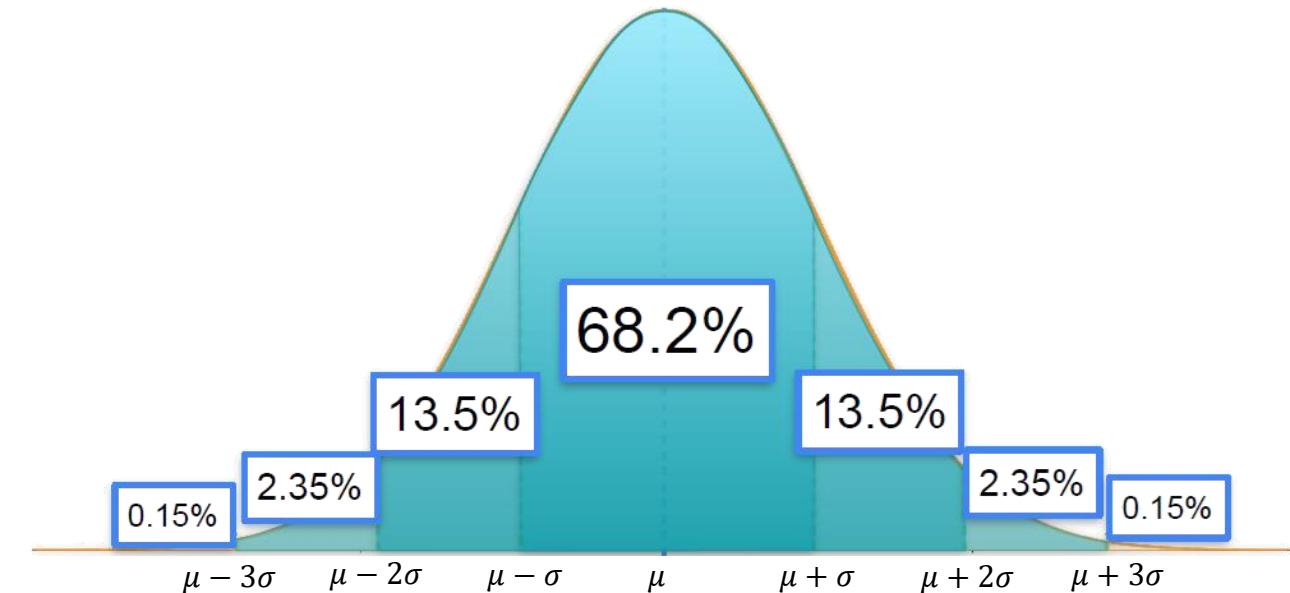
- μ is the mean of the data
- σ is the standard deviation of the data

$$X \sim \mathcal{N}(\mu, \sigma^2)$$



Standard Deviation in the Normal Distribution

- 68-95-99.7
 - μ is the mean of the data
 - σ is the standard deviation of the data
 - $X \sim \mathcal{N}(\mu, \sigma^2)$



...Standard Deviation in the Normal Distribution

- Sum of Normal distributions

- $W = aX + bY$
- X and Y are independent
- $X \sim \mathcal{N}(\mu_X, \sigma_X^2)$
- $Y \sim \mathcal{N}(\mu_Y, \sigma_Y^2)$

$$W \sim \mathcal{N}(a\mu_X + b\mu_Y, a^2\sigma_X^2 + b^2\sigma_Y^2)$$

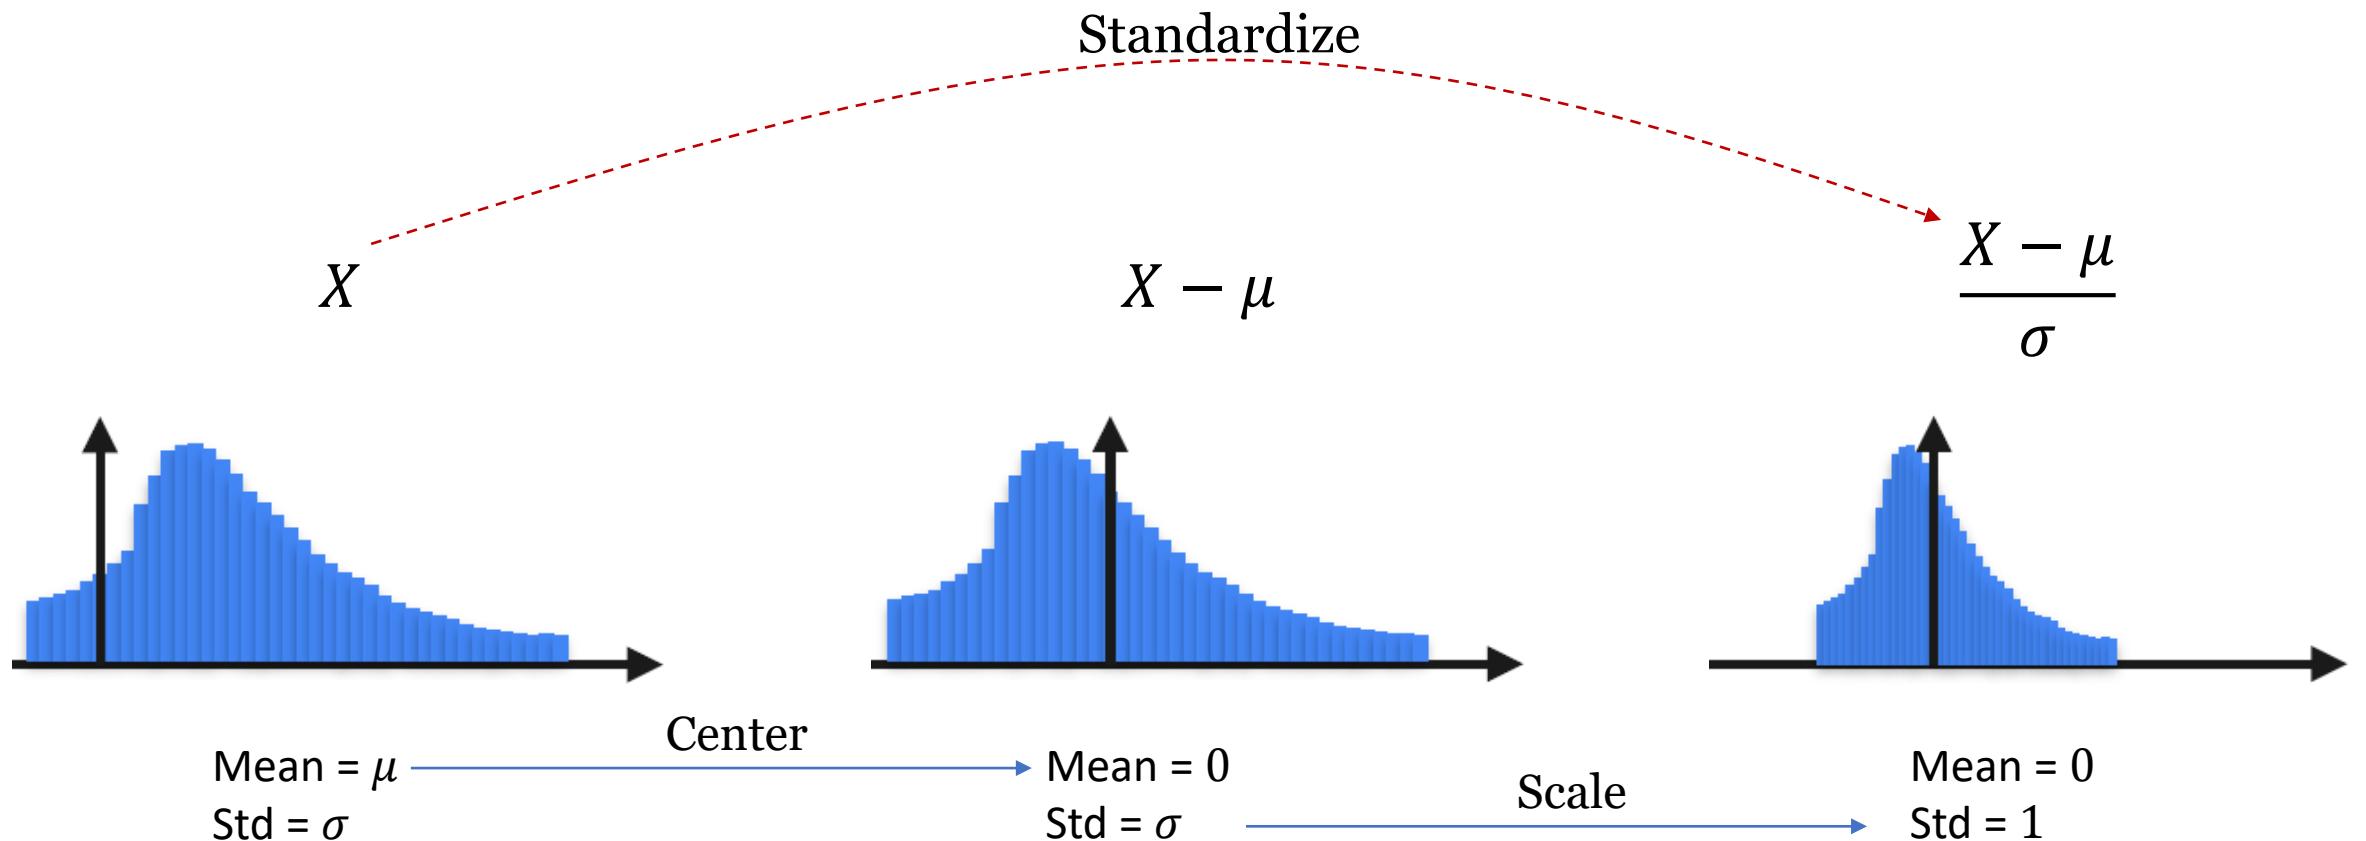


Standardization

- Everything is nicer when the mean is 0 and the standard deviation is 1
- How to standardize data?
 - $X - \mu$ leads to 0 mean
 - $\mathbb{E}[X + Y] = \mathbb{E}[X] + \mathbb{E}[Y] \rightarrow \mathbb{E}[X - \mu] = \mathbb{E}[X] - \mathbb{E}[\mu] = \mathbb{E}[X] - \mu = 0$
 - $\frac{X}{\sigma}$ leads to 1 std
 - $Var(cX) = \mathbb{E}[(cX)^2] - \mathbb{E}[cX]^2 = \mathbb{E}[c^2X^2] - c^2\mathbb{E}[X]^2 = c^2\mathbb{E}[X^2] - c^2\mathbb{E}[X]^2 = c^2(\mathbb{E}[X^2] - \mathbb{E}[X]^2) = c^2Var(X)$
 - $Var\left(\frac{X}{\sigma}\right) = \frac{1}{\sigma^2}Var(X)$
 - $std\left(\frac{X}{\sigma}\right) = \frac{1}{\sigma}std(X) = \frac{\sigma}{\sigma} = 1$



... Standardizing a Distribution



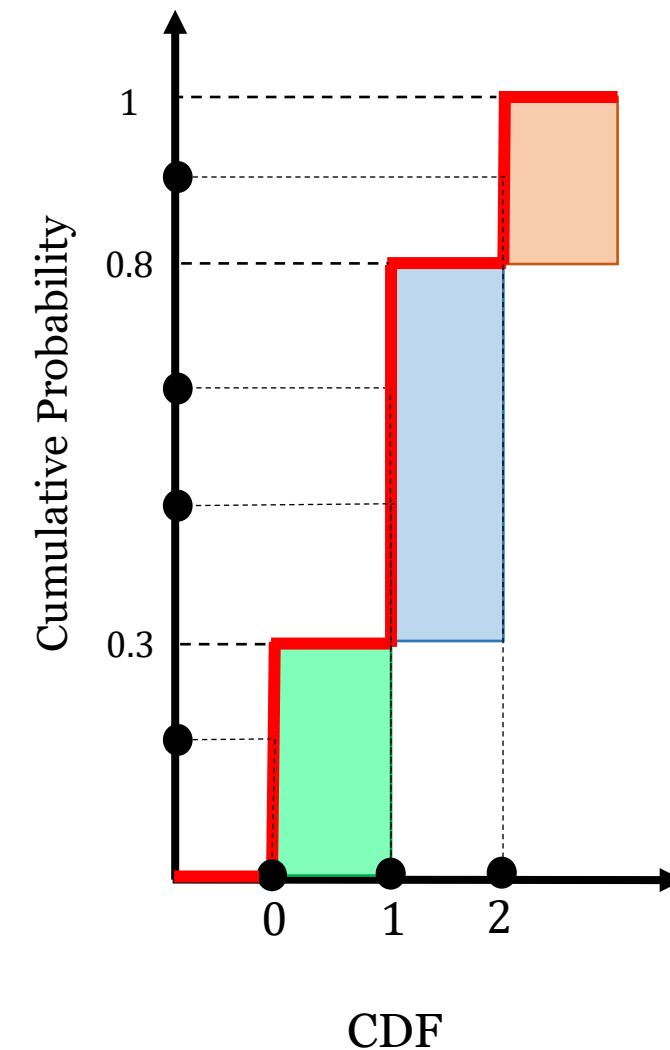
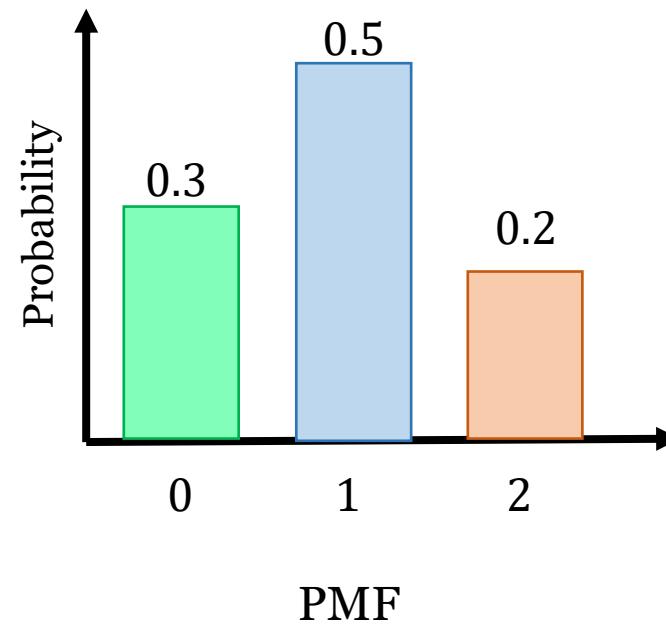
Sampling from a Distribution

- Create syntactic data that looks like the original one
 - We know the original data distribution
 - We need to pick points that have probabilities given by the original distribution
- How?
 - Draw the **cumulative distribution** function
 - Generate numbers **uniformly** between 0 and 1
 - Locate the generated numbers on the vertical axis of CDF
 - Pick related values on the horizontal axis of CDF



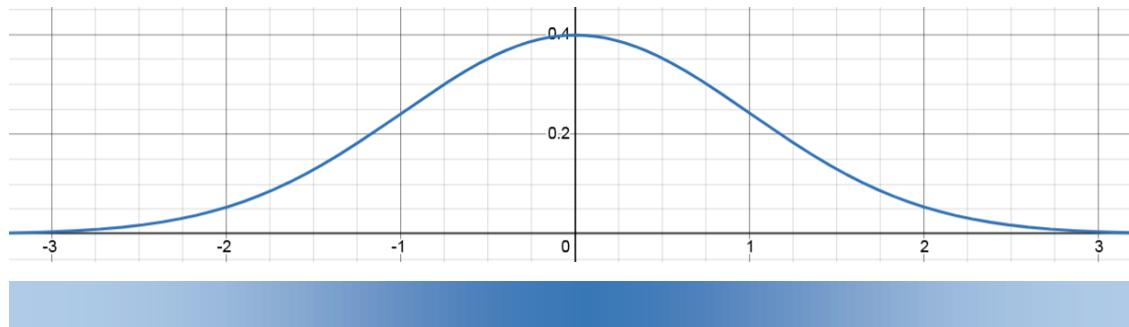
... Sampling from a Distribution

- Discrete distribution

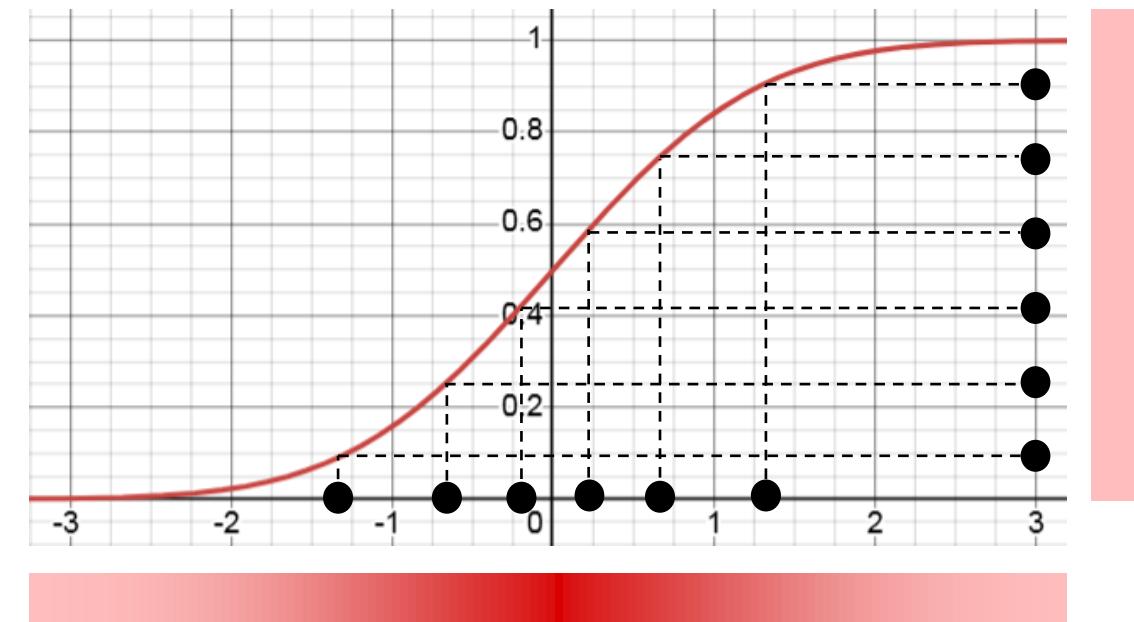


... Sampling from a Distribution

- Continuous distribution



PDF



CDF

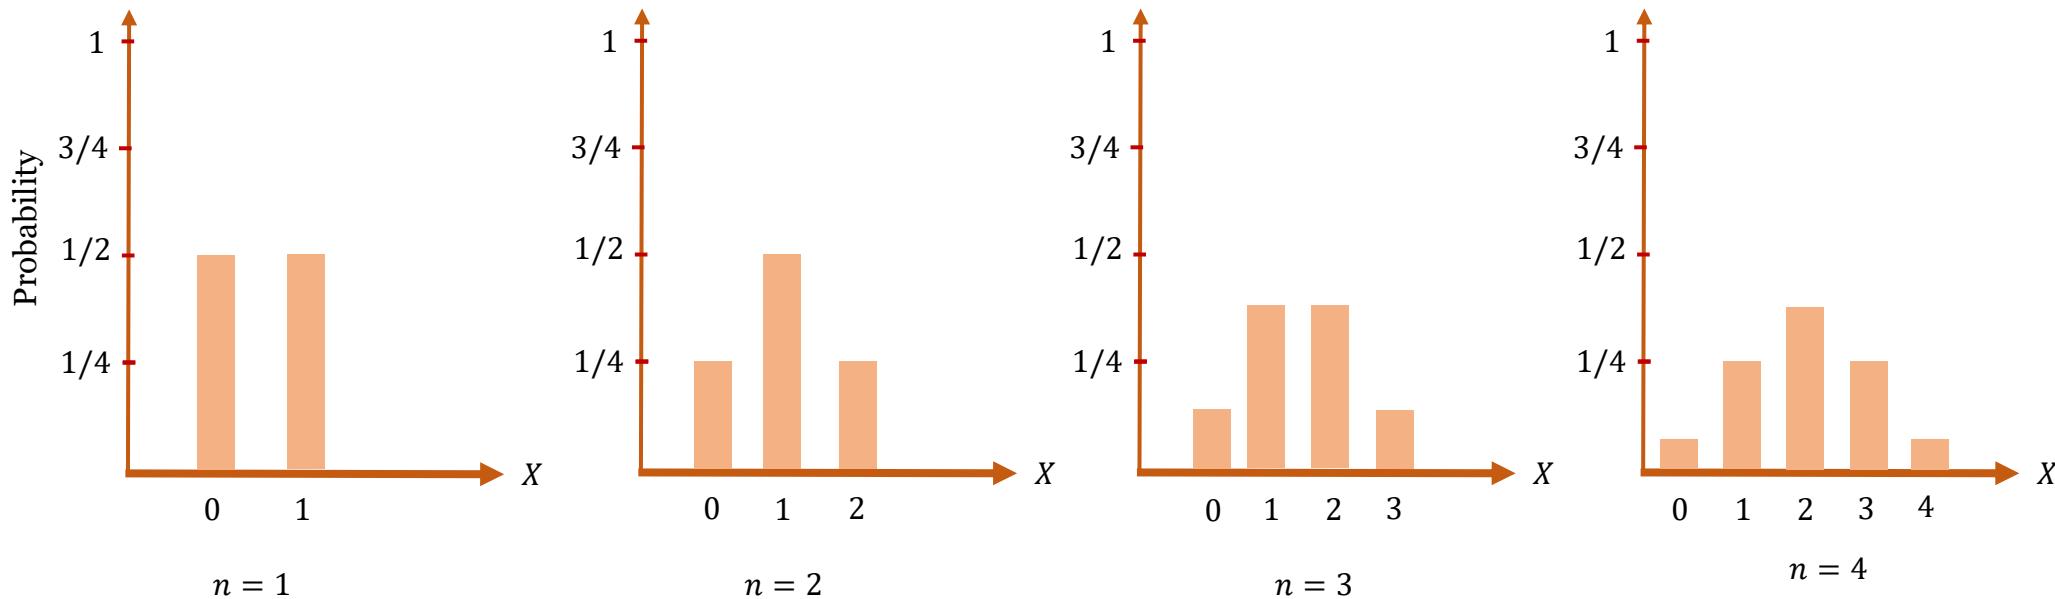


Central Limit Theorem



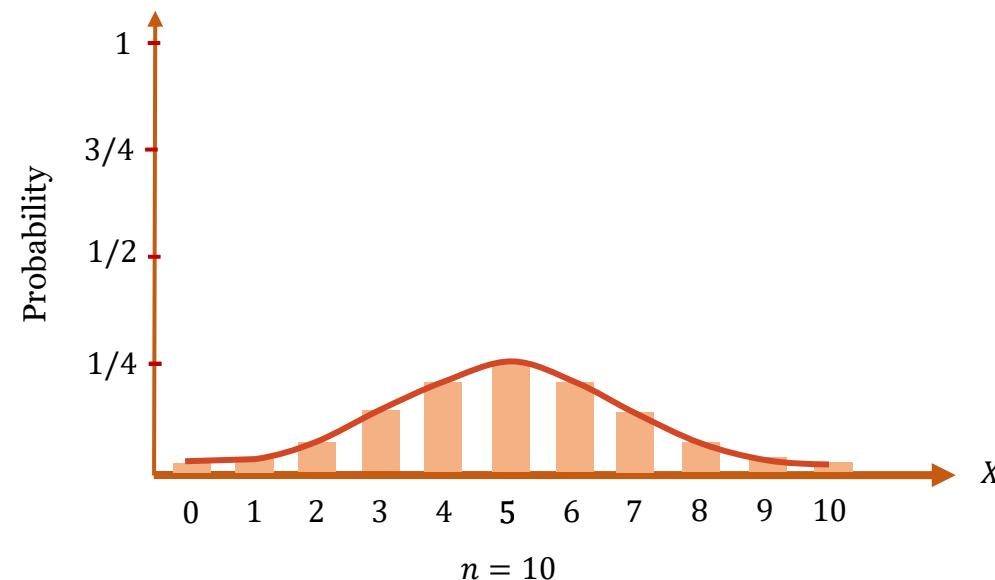
Discrete example

- flip a fair coin n times
 - $P(H) = P(T) = 0.5$
 - Random variable X = number of heads



...Discrete example

- Look like a Gaussian distribution!



As you increase the number of observations, the probability distribution becomes closer to a Gaussian distribution.

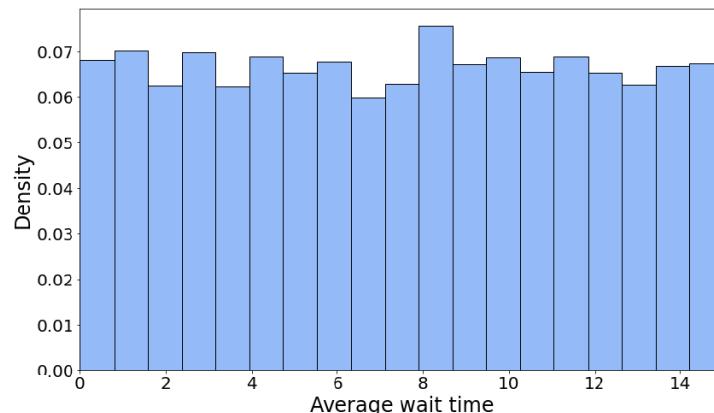


Continuous example

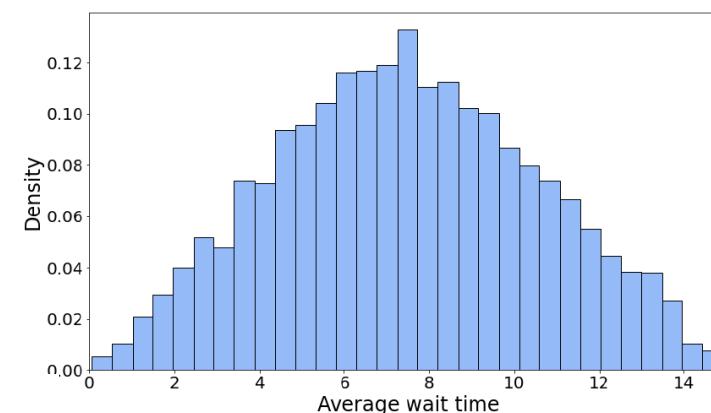
- Random variable X = wait to time for a call to be answered
- Assume that X follows a uniform distribution: $X \sim \mathcal{U}(0, 15)$
- Calculate average waiting time (Y) using n samples
 - $n = 1 \rightarrow Y_1 = \frac{X_1}{1}$
 - $n = 2 \rightarrow Y_2 = \frac{X_1 + X_2}{2}$
 - ...
 - $Y_n = \frac{1}{n} \sum_{i=1}^n X_i$
- What is the distribution of Y_n
 - Again, Gaussian!



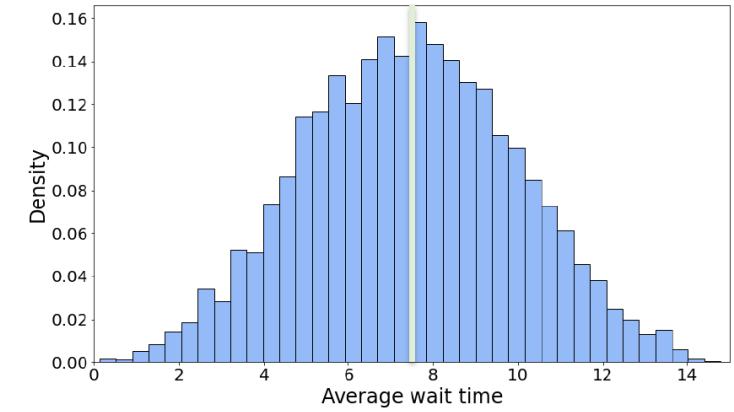
...Continuous example



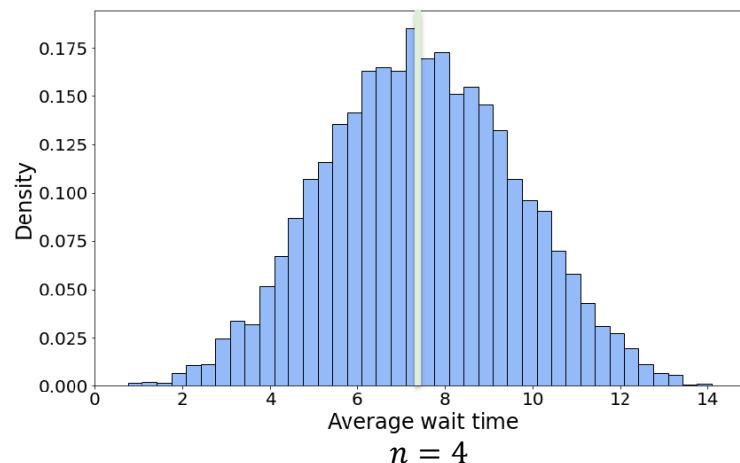
$n = 1$



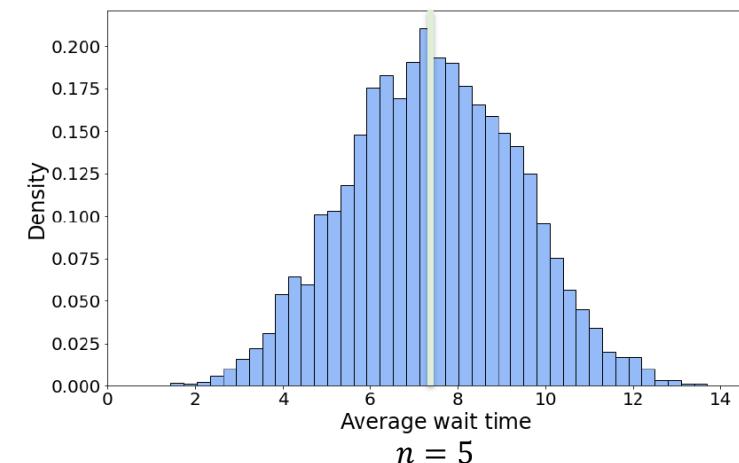
$n = 2$



$n = 3$

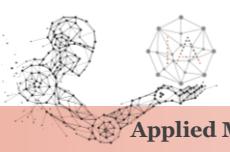


$n = 4$



$n = 5$

When you average on a large enough number of samples, the distribution will approximately follow a normal distribution



...Continuous example

- Mean & variance of Y_n
 - $\mathbb{E}[Y_n] = \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n X_i\right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] = \frac{1}{n} n \mathbb{E}[X] = \mathbb{E}[X]$
 - $\mu_{Y_n} = \mu$
 - Mean of Y_n equals the population mean
- $Var(Y_n) = Var\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n^2} \sum_{i=1}^n Var(X_i) = \frac{1}{n^2} n Var(X) = \frac{Var(X)}{n}$
 - $\sigma_{Y_n}^2 = \frac{\sigma^2}{n}$
 - Variance of Y_n equals $\frac{1}{n}$ of the population variance



Central Limit Theorem (Definition)

- Formal definition
 - Regardless of the distribution of X , as $n \rightarrow \infty$, the variable $\frac{1}{n} \sum_{i=1}^n X_i$ follows a normal distribution. We can standardize it and get the standard normal distribution.
 - $\frac{Y_n - \mu_{Y_n}}{\sigma_Y} \sim \mathcal{N}(0, 1^2)$
 - $$\begin{aligned} \frac{Y_n - \mu_{Y_n}}{\sigma_Y} &= \frac{\frac{1}{n} \sum_{i=1}^n X_i - \mathbb{E}[X]}{\sqrt{\frac{\sigma_X^2}{n}}} = \frac{\frac{1}{n} \sum_{i=1}^n X_i - \mathbb{E}[X]}{\sigma_X} \sqrt{n} = \frac{\frac{1}{n} \sum_{i=1}^n X_i - \frac{1}{n} n \mathbb{E}[X]}{\sigma_X} \sqrt{n} = \frac{\frac{1}{n} \sum_{i=1}^n X_i - n \mathbb{E}[X]}{\sigma_X} \sqrt{n} \\ &= \frac{\sum_{i=1}^n X_i - n \mathbb{E}[X]}{\sqrt{n} \sigma_X} \end{aligned}$$
 - $\frac{\sum_{i=1}^n X_i - n \mathbb{E}[X]}{\sqrt{n} \sigma_X} \sim \mathcal{N}(0, 1^2)$

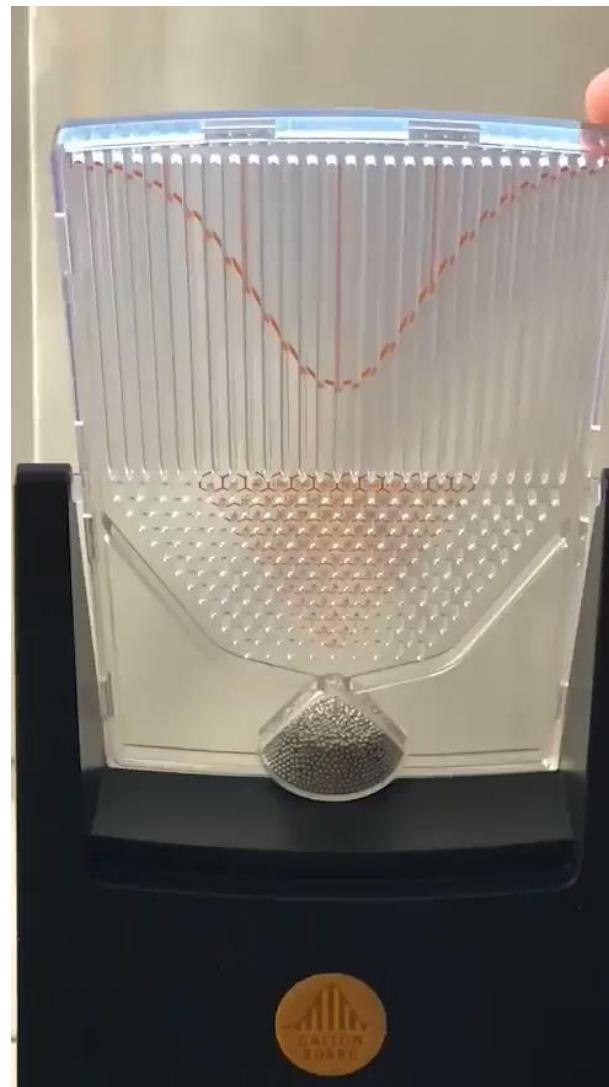


...Central Limit Theorem (Definition)

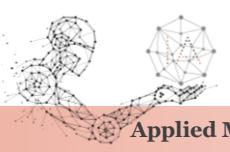
- Narrative description
 - Take **any arbitrary** distribution, take a few samples (always the same number), and look at the average
 - Do this many times and plot all these averages. You get the normal distribution.
 - No matter what distribution you started with in the first place.
- Note
 - In general, a safe rule is that you usually need about 30 samples before the bell-shaped distribution comes in.
 - If the original population of the data is very skewed, you usually need more samples than if you are working with a symmetric distribution.



Galton board



Multiple Variables

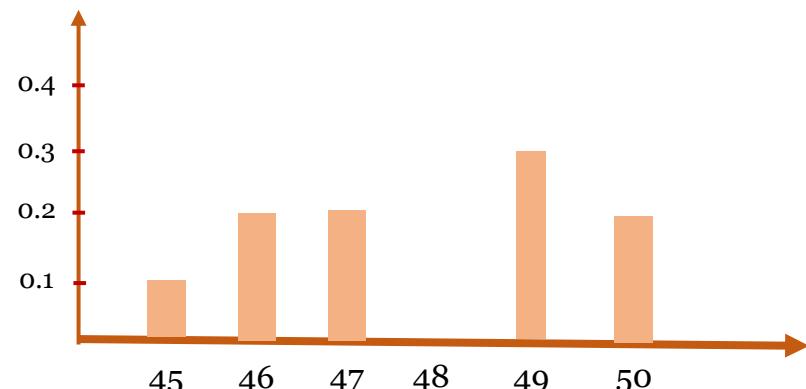
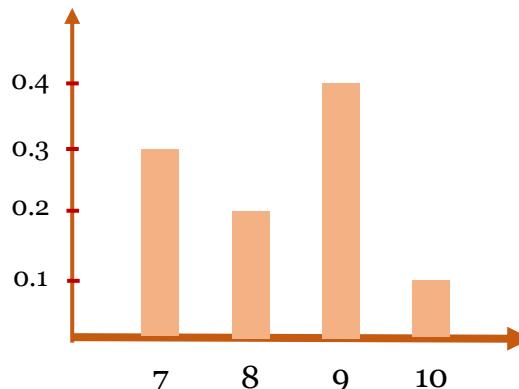


Joint Distributions (Discrete)

- Example-1

Age	Count
7	3
8	2
9	4
10	1

Height	Count
45	1
46	2
47	2
48	0
49	3
50	2



... Joint Distributions (Discrete)

- ... Example-1

Age	Height
7	45
7	46
7	46
8	47
8	47
9	49
9	49
9	49
9	50
10	50



	45	46	47	48	49	50
Age						
7	1	2	0	0	0	0
8	0	0	2	0	0	0
9	0	0	0	0	3	1
10	0	0	0	0	0	1



... Joint Distributions (Discrete)

- ... Example-1

		Height					
		45	46	47	48	49	50
Age	7	1	2	0	0	0	0
	8	0	0	2	0	0	0
	9	0	0	0	0	3	0
	10	0	0	0	0	0	0

Divide by sum

		Height (Y)					
		45	46	47	48	49	50
Age (X)	7	1/10	2/10	0	0	0	0
	8	0	0	2/10	0	0	0
	9	0	0	0	0	3/10	1/10
	10	0	0	0	0	0	1/10

Probability Mass Function

All probabilities for all possible combinations of X and Y

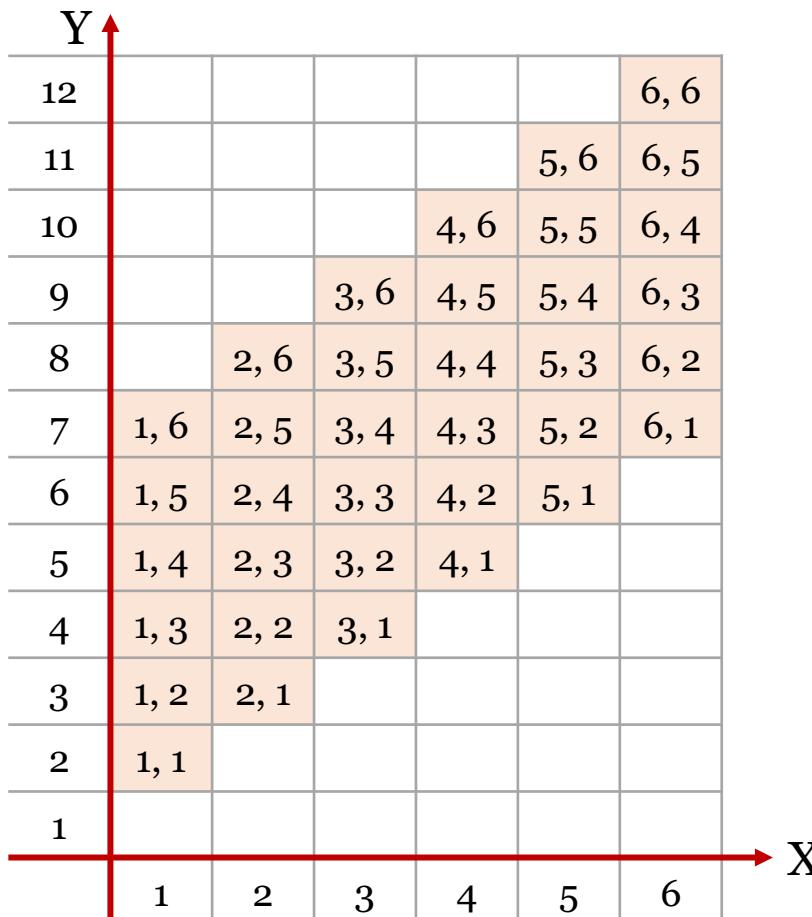
Joint Distribution

$$p_{XY}(x, y) = P(X = x, Y = y)$$



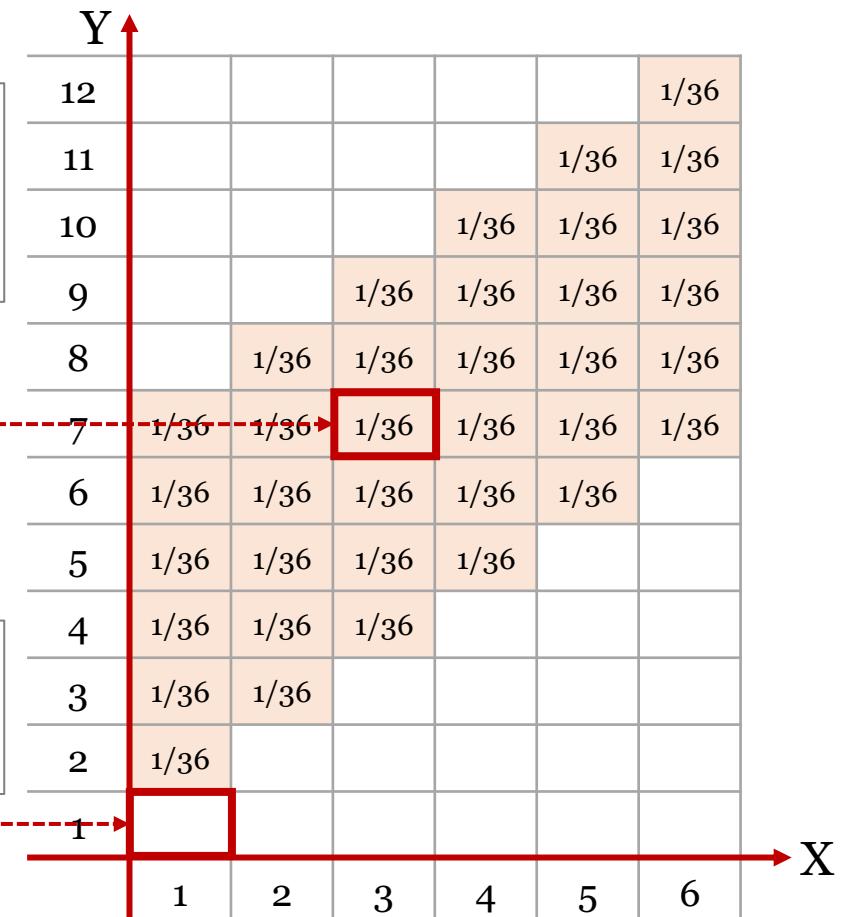
... Joint Distributions (Discrete)

- Example-2- Toss two dices X = the number rolled on the 1st dice Y = sum of the two dices



Example:
 $p_{XY}(3,7)$
 $= P(X = 3, Y = 7) = \frac{1}{36}$

Example:
 $p_{XY}(1,1)$
 $= P(X = 1, Y = 1) = 0$



... Joint Distributions (Discrete)

- For independent discrete variables
 - $p_{XY}(x, y) = P(X = x, Y = y) = P(x) \times P(y)$
- Example- Toss two dices
 - X = the number rolled on the 1st dice
 - Y = the number rolled on the 2nd dice
 - $p_{XY}(2,3) = P(2) \times P(3) = \frac{1}{6} \times \frac{1}{6} = \frac{1}{36}$



Marginal Distribution

- Distribution of one variable while ignoring the other
- To find the marginal distribution for a variable, sum the joint probability distribution over all value of the other variable.
- Formula

$$P_Y(y_j) = \sum_i P_{XY}(x_i, y_j)$$



... Marginal Distribution

- Example
 - Age and height dataset

Age	7	7	7	8	8	9	9	9	9	10
Height	45	46	46	47	47	49	49	49	50	50

- Joint distribution

		Height (Y)					
		45	46	47	48	49	50
Age (X)	7	1/10	2/10	0	0	0	0
	8	0	0	2/10	0	0	0
	9	0	0	0	0	3/10	1/10
	10	0	0	0	0	0	1/10



... Marginal Distribution

- ... Example
 - Marginal distribution of height → Do NOT care about age anymore and only care about height
 - Summarize the behavior of the distribution across only the height variable
 - Add over each height value all the probabilities

		Height (Y)							Marginal distribution of age
		45	46	47	48	49	50		
Age (X)	7	1/10	2/10	0	0	0	0	3/10	↓
	8	0	0	2/10	0	0	0	2/10	
	9	0	0	0	0	3/10	1/10	4/10	
	10	0	0	0	0	0	1/10	1/10	

Marginal distribution of height → $\begin{matrix} 1/10 & 2/10 & 2/10 & 0 & 3/10 & 2/10 \end{matrix}$



Conditional Distribution

- Example
 - Age and height dataset
 - Only care about age 9
 - If age = 9, what is the distribution across the height variable?
 - $P_{Y|X=9}(y) = P(Y = y|X = 9)$

		Height (Y)					
		45	46	47	48	49	50
Age (X)	7	1/10	2/10	0	0	0	0
	8	0	0	2/10	0	0	0
	9	0	0	0	0	3/10	1/10
	10	0	0	0	0	0	1/10

Example:

$$p_{Y|X=9}(49) = P(Y = 49|X = 9) = \frac{3}{4}$$

		Height (Y)					
		45	46	47	48	49	50
Age (X)	9	0	0	0	0	3/10	1/10

Sum of the probabilities does not equal 1!

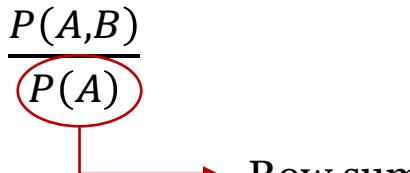
Divide by row sum Normalize

9	0	0	0	0	3/4	1/4
---	---	---	---	---	-----	-----

... Conditional Distribution

- Observe one variable given that the value of the other one is known
- Why divide by the row sum?

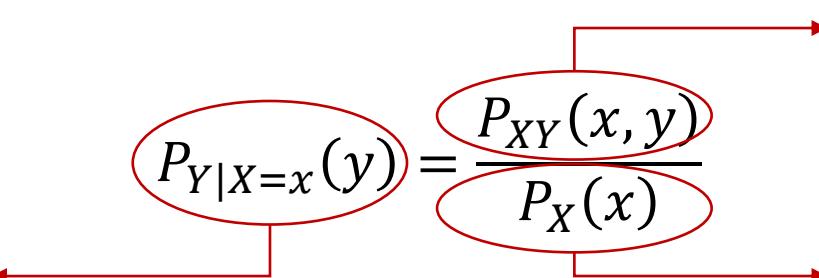
- $P(A, B) = P(A) \times P(B|A) \Rightarrow P(B|A) = \frac{P(A,B)}{P(A)}$



- Example

- $P_{Y|X=9}(y) = P(Y = y|X = 9) = \frac{P(X=9,Y=y)}{P(X=9)}$
- $P_{Y|X=9}(49) = P(Y = 49|X = 9) = \frac{P(X=9,Y=49)}{P(X=9)} = \frac{\frac{3}{10}}{\frac{4}{10}} = \frac{3}{4}$

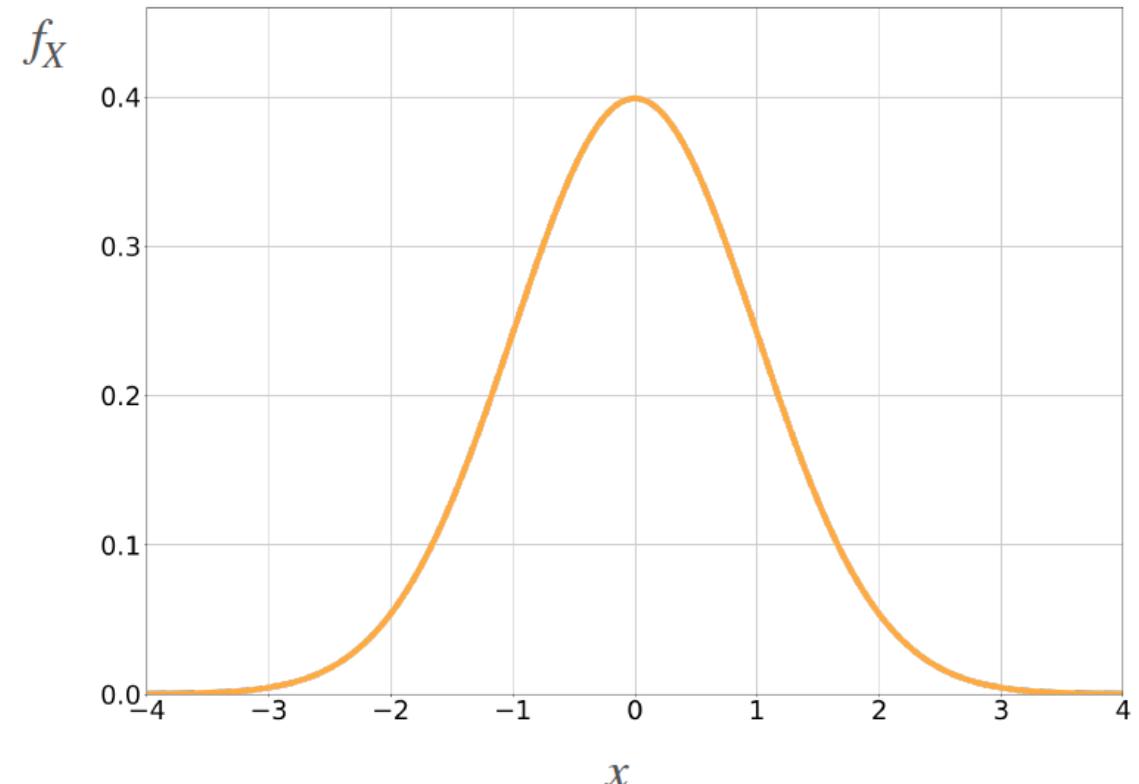
- General formula

$$P_{Y|X=x}(y) = \frac{P_{XY}(x, y)}{P_X(x)}$$




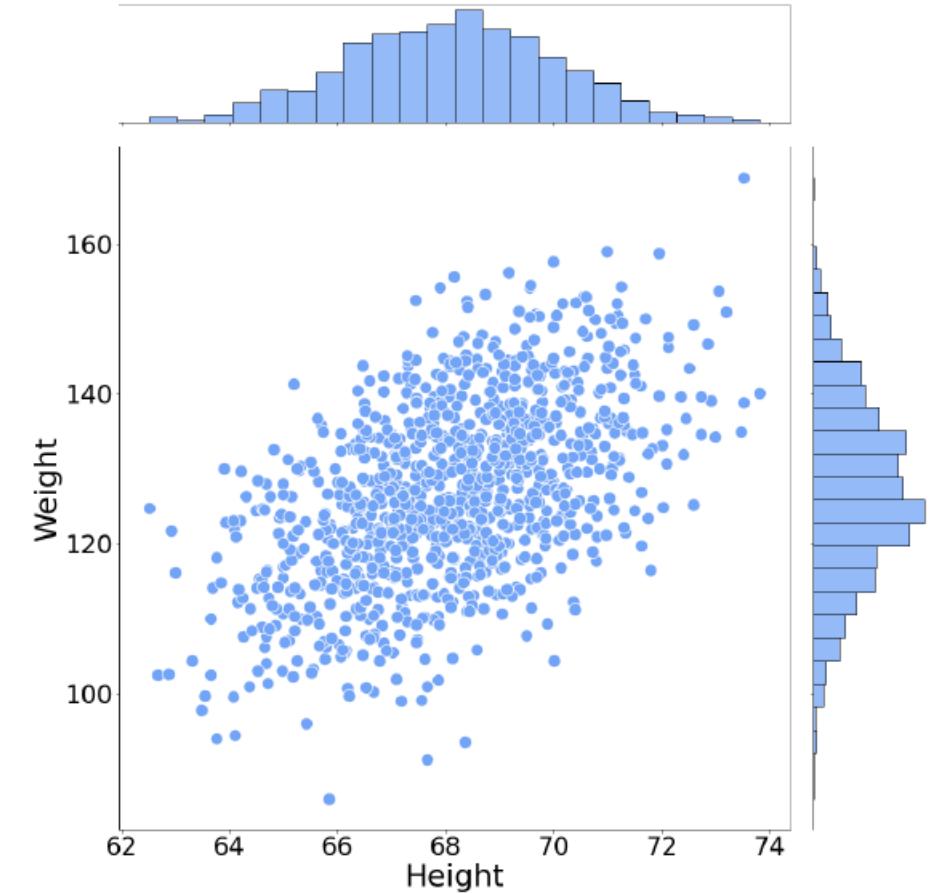
Multivariate Gaussian Distribution

- For a single variable
 - $X \sim \mathcal{N}(\mu, \sigma^2)$
 - μ is the mean of the data
 - σ is the standard deviation of the data
 - $f_X(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$
- What if we have more than one variable?



... Multivariate Gaussian Distribution

- Example
 - H: Height of an adult in inches
 - $H \sim \mathcal{N}(\mu_H, \sigma_H^2)$
 - W: Weight of an adult in pounds
 - $W \sim \mathcal{N}(\mu_W, \sigma_W^2)$
 - If H and W were independent variables we could calculate $f_{HW}(h, w)$ simply by multiplying $f_H(h)$ and $f_W(w)$
 - But it is not the case here
 - Taller adults tend to be fatter
 - There is a correlation between variables



... Multivariate Gaussian Distribution

- General formula is very similar to the formula of the single one ($\frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$)
 - Use covariance matrix instead of the variance
 - For vectors instead of scalar values

$$x = [x_1 \quad x_2 \quad \dots \quad x_n]^T$$
$$\mu = [\mu_1 \quad \mu_2 \quad \dots \quad \mu_n]^T$$
$$f(x_1, x_2, \dots, x_n) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1} (x-\mu)}$$



Confidence Interval



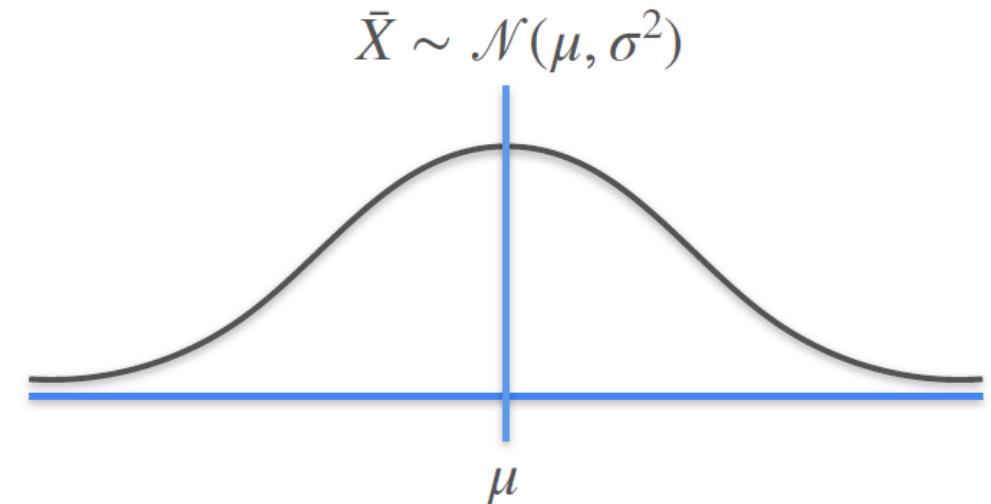
Introduction

- Example
 - We can use sampling to estimate the mean of population height
 - How good our estimate is?
 - How can we assure the estimation is close to the actual population mean?
- Confidence interval is a technique that is used to establish a **degree of certainty** to the point estimates from the samples
- Example
 - Instead of \bar{x} , we use a lower and upper bounds to report a range to use the sample mean with some degree of certainty
 - $lower\ limit < \bar{x} < upper\ limit$



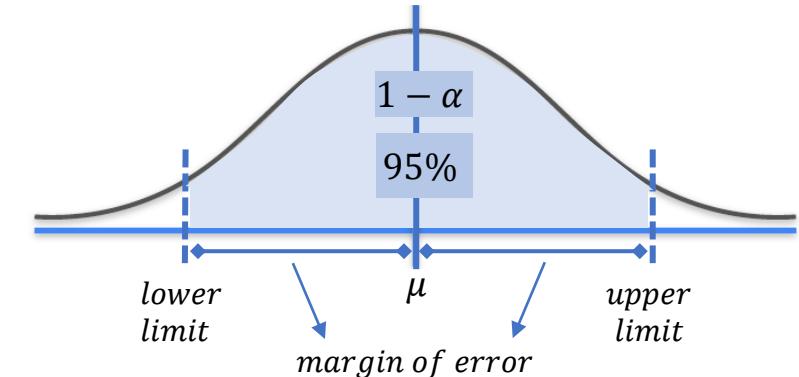
... Introduction

- Example
 - Sample size = 1
 - Find the mean of the sample to use as the estimate for the population mean
 - We take multiple samples of size 1
 - The sampling distribution for the sample means (\bar{X}) is a Gaussian distribution $\mathcal{N}(\mu, \sigma^2)$
 - Based on the Central Limit Theorem
 - μ = population mean
 - σ = population standard deviation



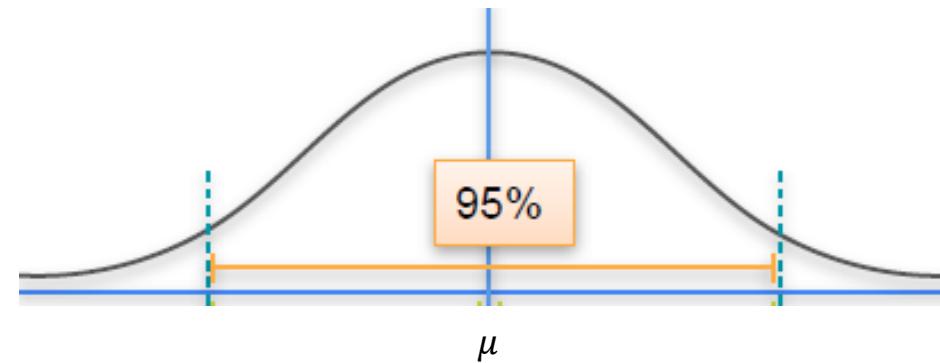
... Introduction

- We want to indicate the frequency with which our sample means lie within an interval (between an upper and a lower limit). How?
 - Decide on a value called the **significance level (α)**
 - Common value is **0.05**
 - The value **$1 - \alpha$** is called the **confidence level**
 - Common value is **0.95** (or 95%)
 - The confidence level indicates the region that contains 95% of the entire sample means in the sampling distribution
 - The lower and upper limits can be calculated based on the standard deviation and the confidence interval we have chosen.
 - This is called the **confidence interval**
 - The **margin error** is the distance between the upper/lower limits and the true mean



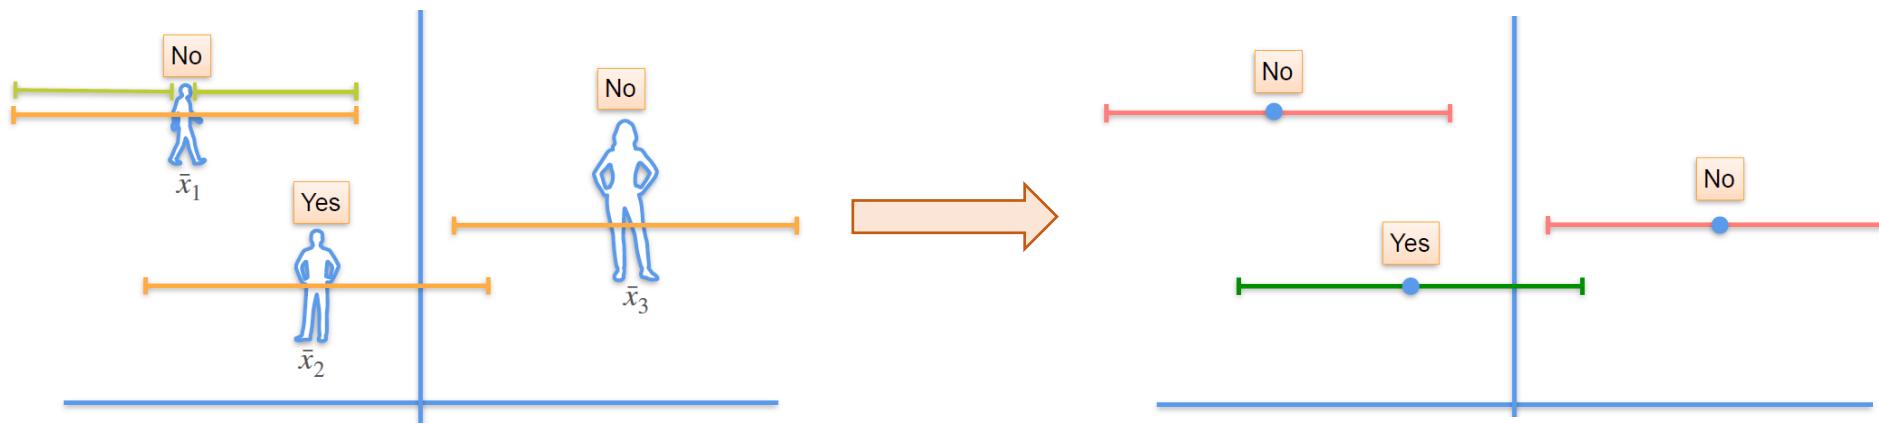
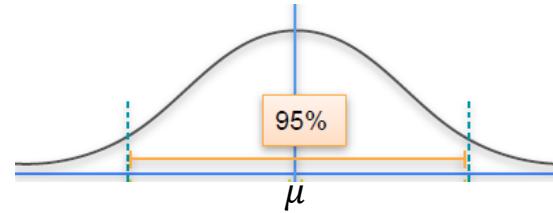
... Introduction

- ... Example- sample size = 1
 - Assume that we know σ
 - We calculate the confidence interval for the known confidence level (95%)



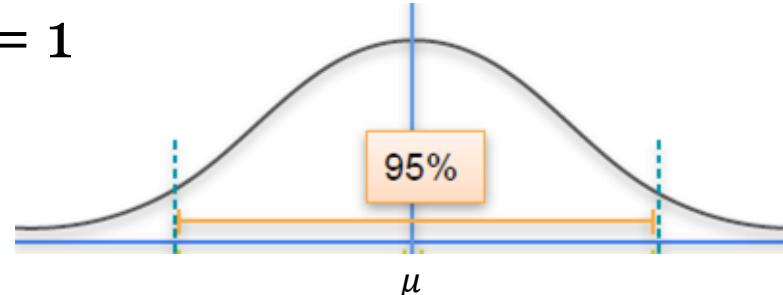
... Introduction

- ... Example- sample size = 1

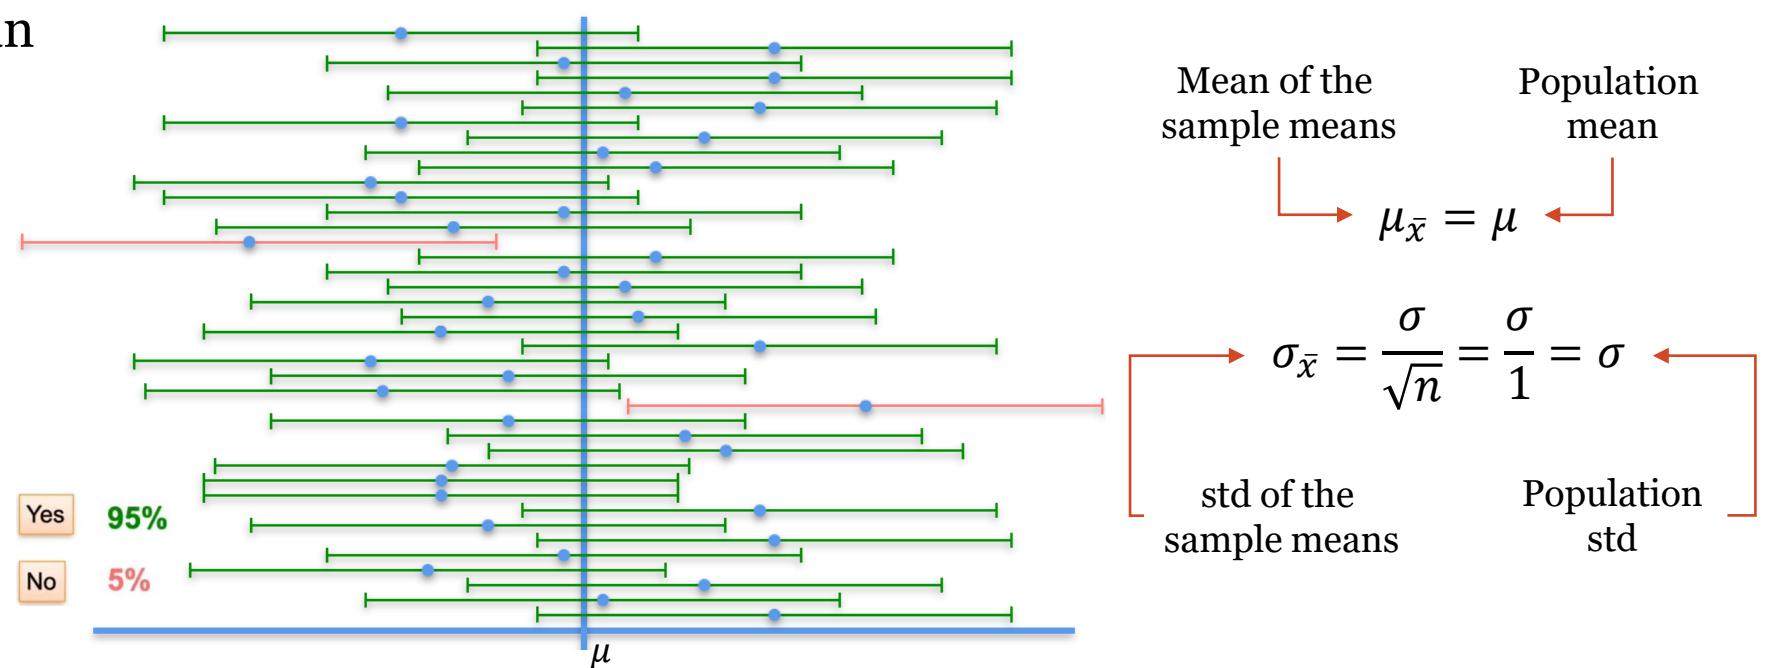


... Introduction

- ... Example- sample size = 1



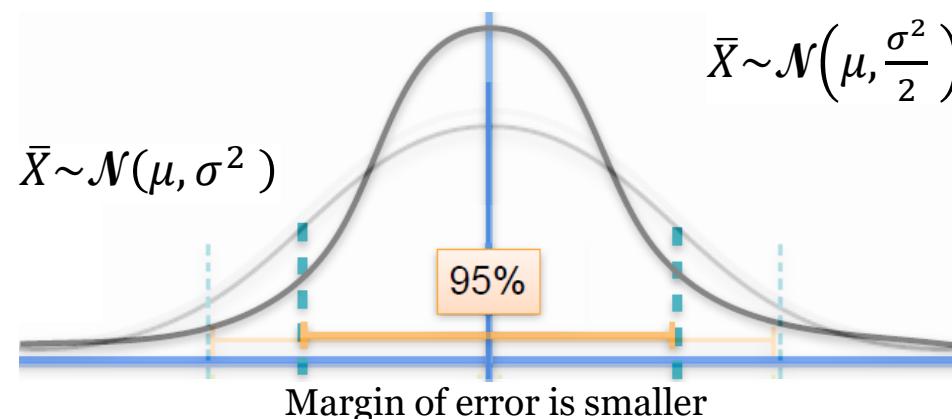
- Repeat for 100 times
 - 95% of the time, the confidence intervals around the samples (mean) contain the population mean



... Introduction

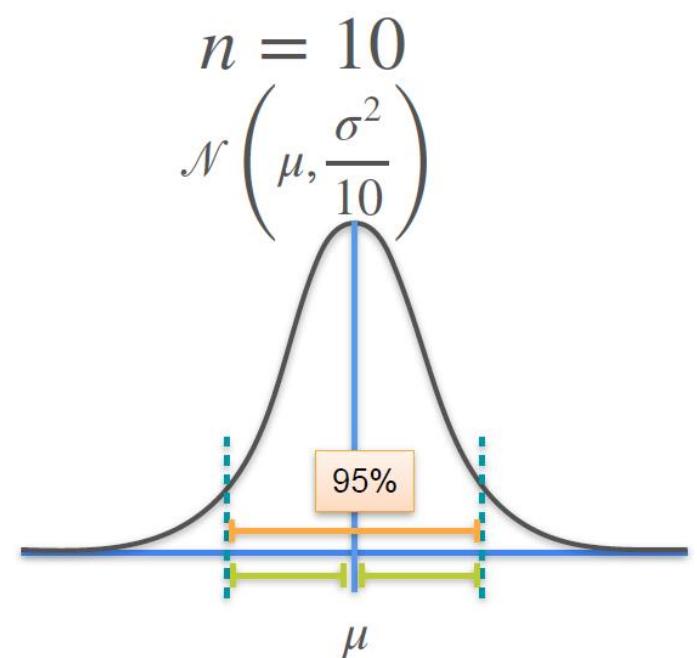
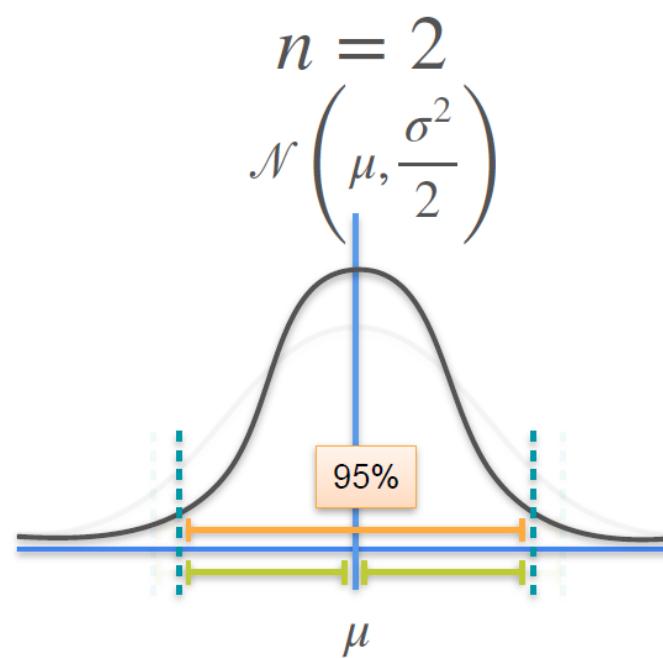
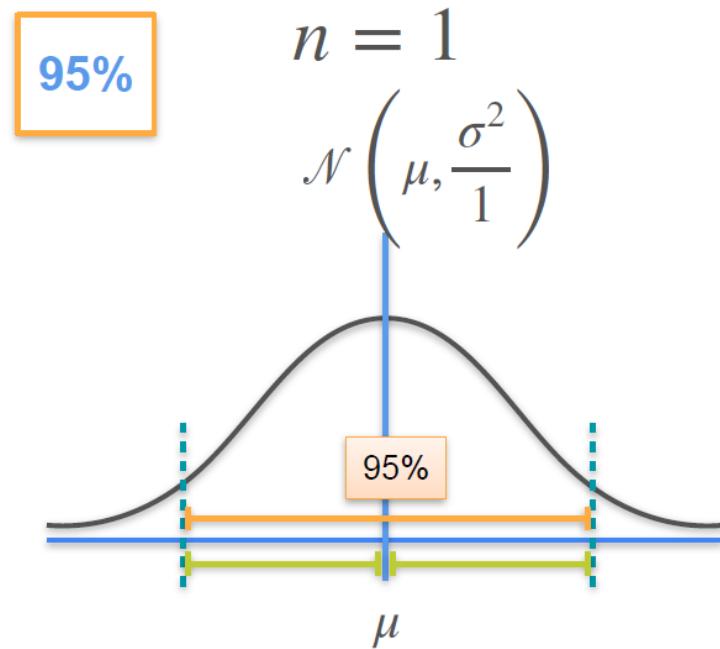
- ... Example- sample size = 2

$$\begin{array}{ccc} \text{Mean of the} & & \text{Population} \\ \text{sample means} & & \text{mean} \\ \text{---} \rightarrow \mu_{\bar{x}} = \mu & & \leftarrow \text{---} \end{array}$$
$$\begin{array}{ccc} \text{---} \rightarrow \sigma_{\bar{x}} = \frac{\sigma}{\sqrt{n}} = \frac{\sigma}{\sqrt{2}} & & \leftarrow \text{---} \\ & & \\ \text{std of the} & & \text{Population} \\ \text{sample means} & & \text{std} \end{array}$$



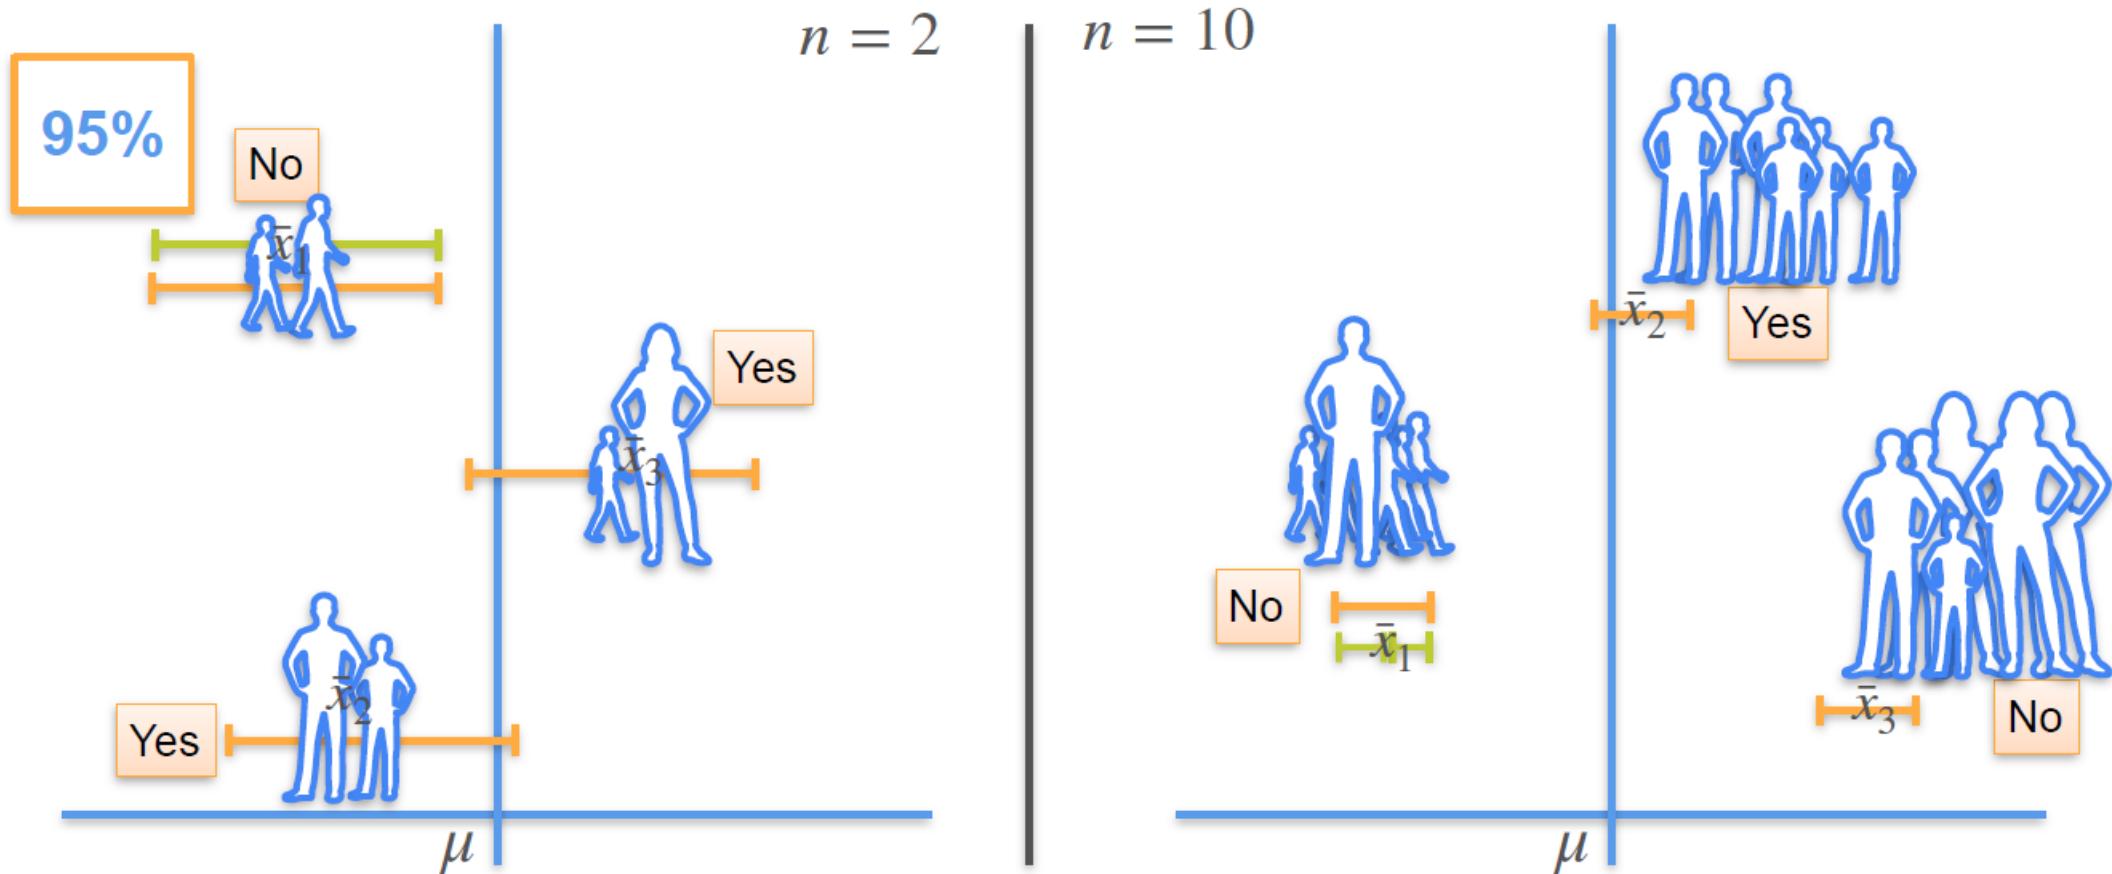
... Introduction

- Effect of the sample size



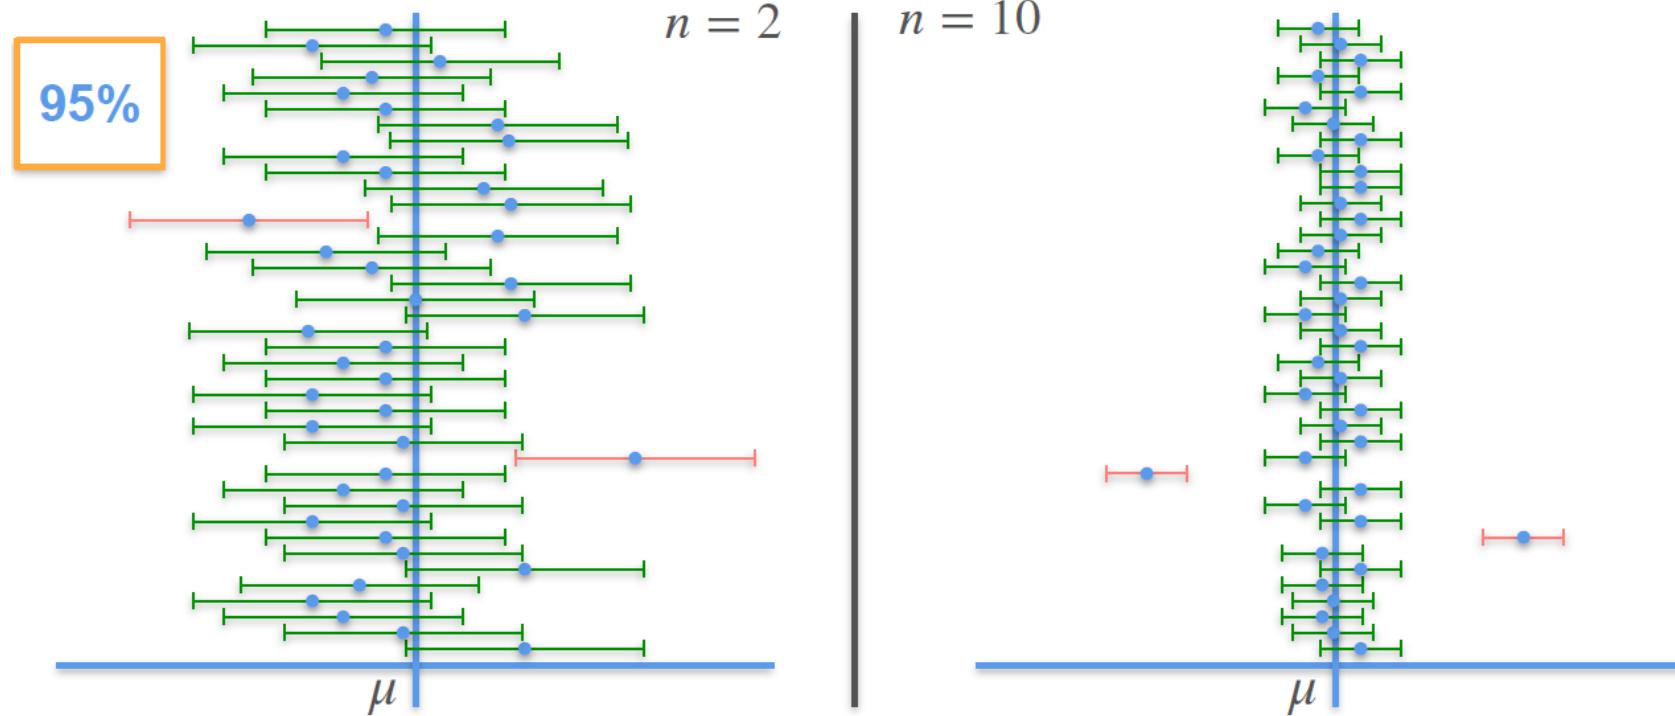
... Introduction

- ... Effect of the sample size



... Introduction

- ... Effect of the sample size



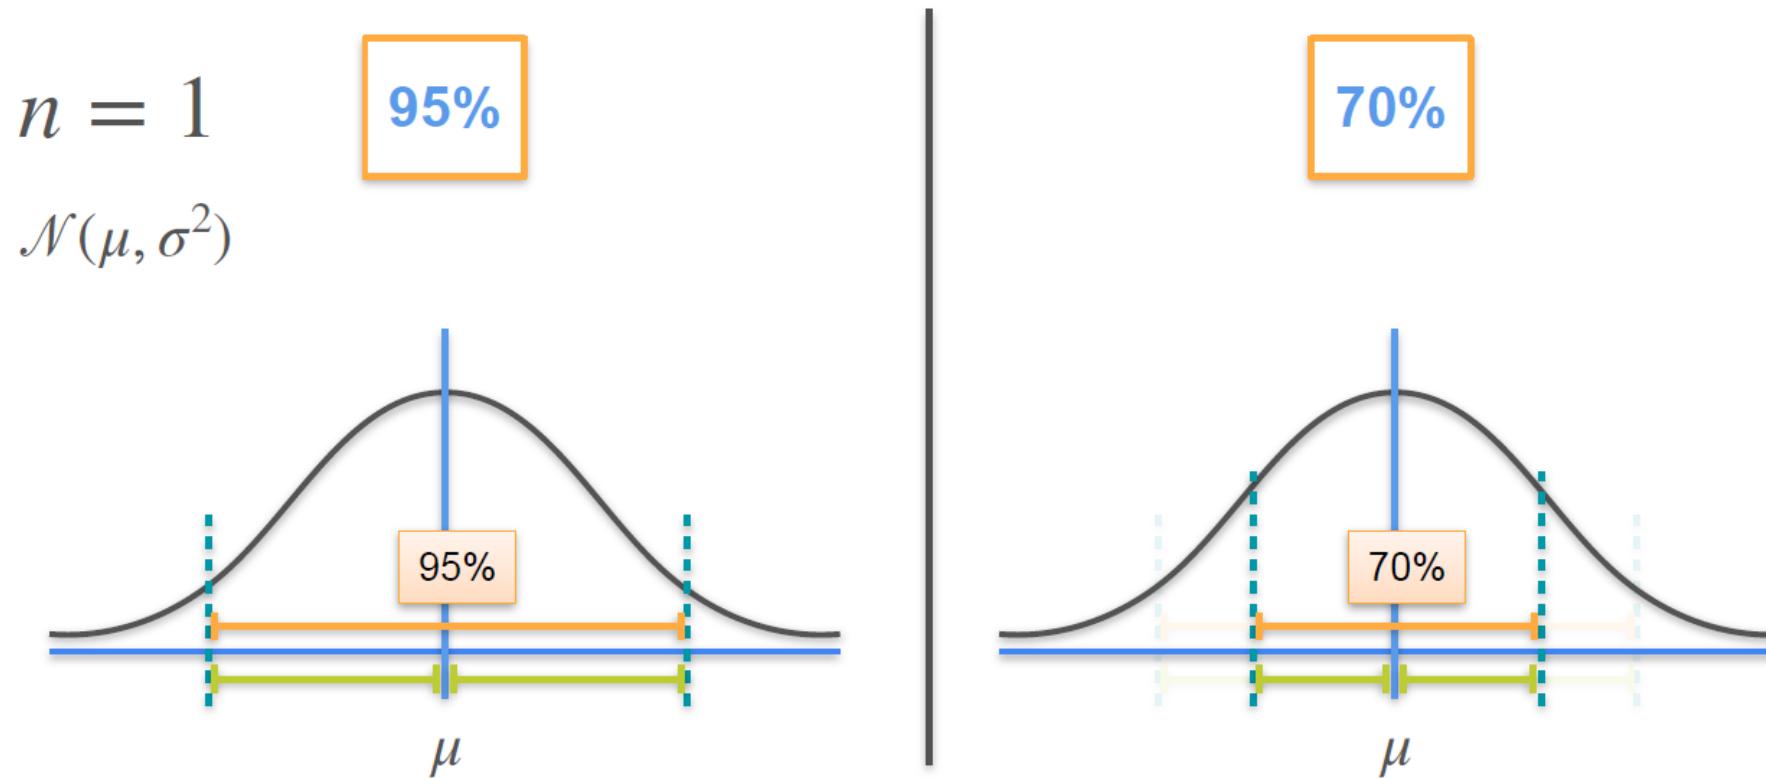
Although we have smaller intervals, we still hit the population mean 95% of the time. **Why?**

Because the sample means are going to be much closer to the population mean



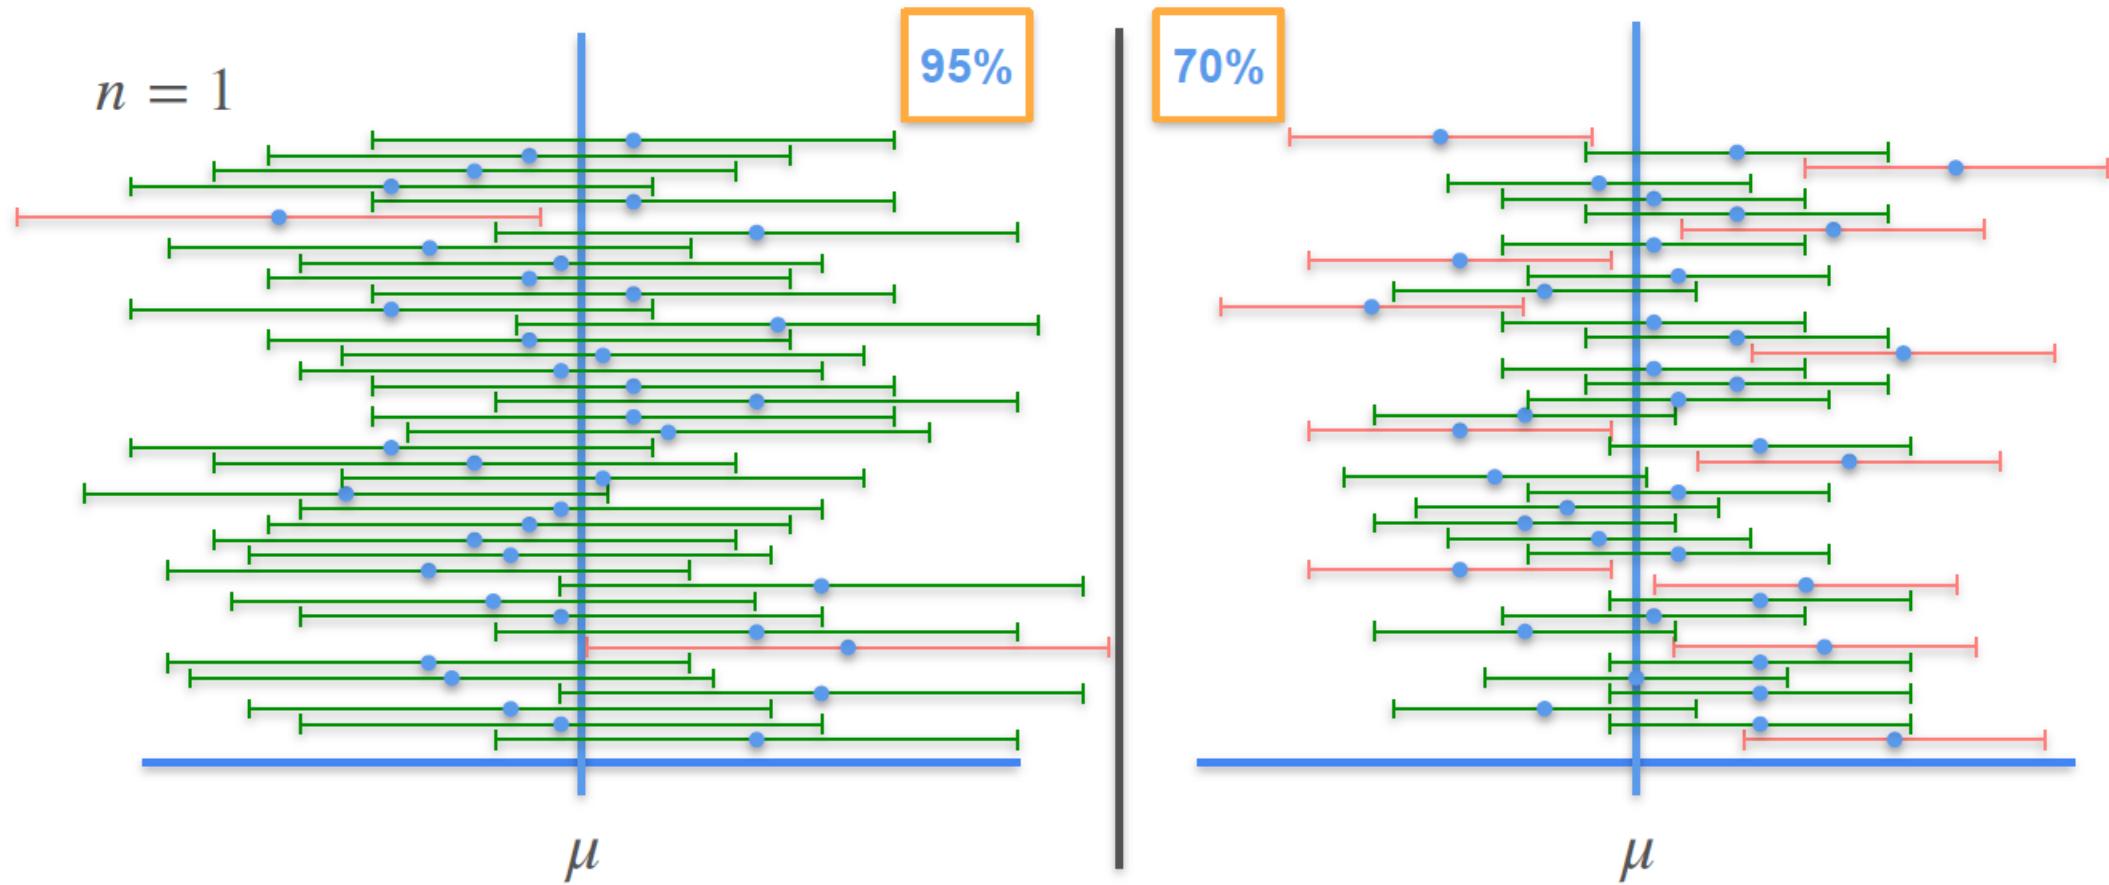
... Introduction

- Effect of the confidence level



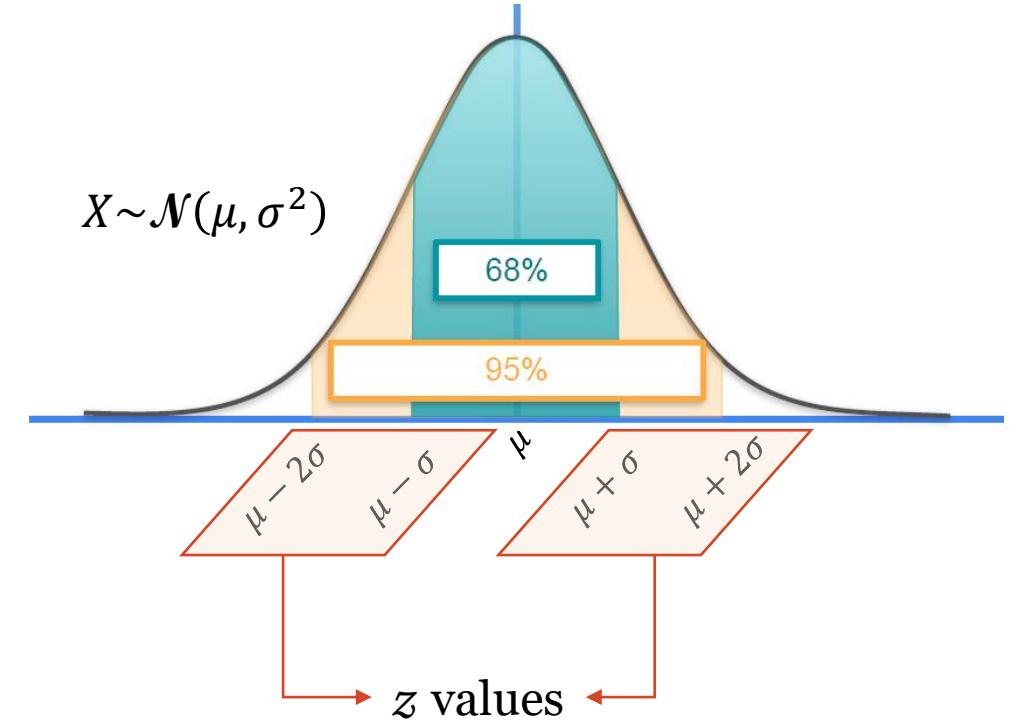
... Introduction

- ... Effect of the confidence level



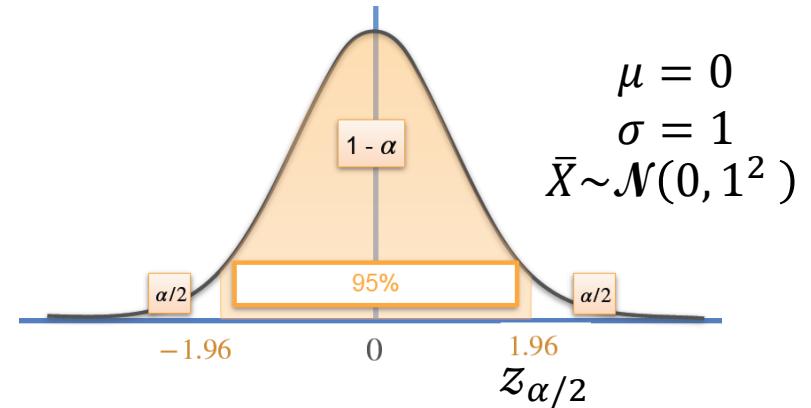
Confidence Interval Calculation

- Margin of error
 - Recall from previous slides
 - 95% of the population falls $2 \times \sigma$ away from the mean
- Standardize the Normal distribution
 - $\bar{X} \sim \mathcal{N}(0,1)$
 - $\mu \pm 2\sigma \sim \pm 2$
 - The exact points ± 1.96

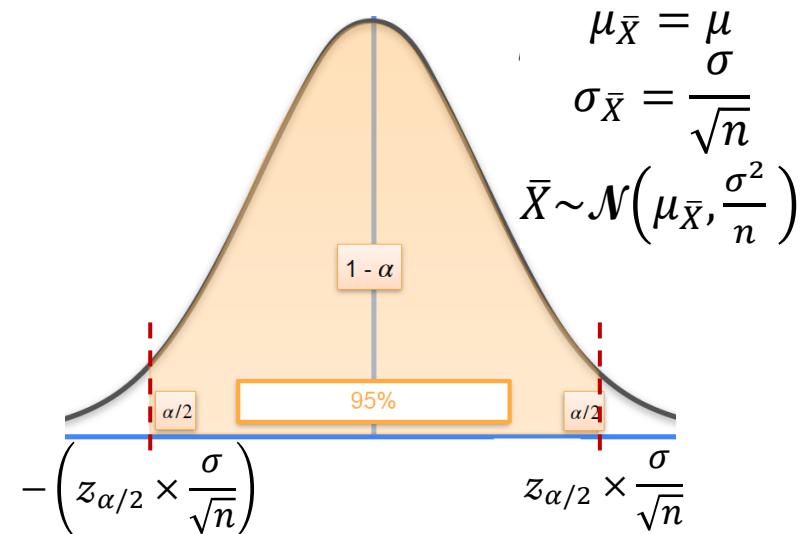


... Confidence Interval Calculation

- ... Margin of error
 - For the standardized Normal distribution, $z_{\alpha/2}$ is called the **critical value**
 - For 95% confidence level $z_{\alpha/2}$ is 1.96



- $\sigma_{\bar{x}} = \frac{\sigma}{\sqrt{n}}$ is called the **standard error**
- Margin of error = $z_{\alpha/2} \times \frac{\sigma}{\sqrt{n}}$

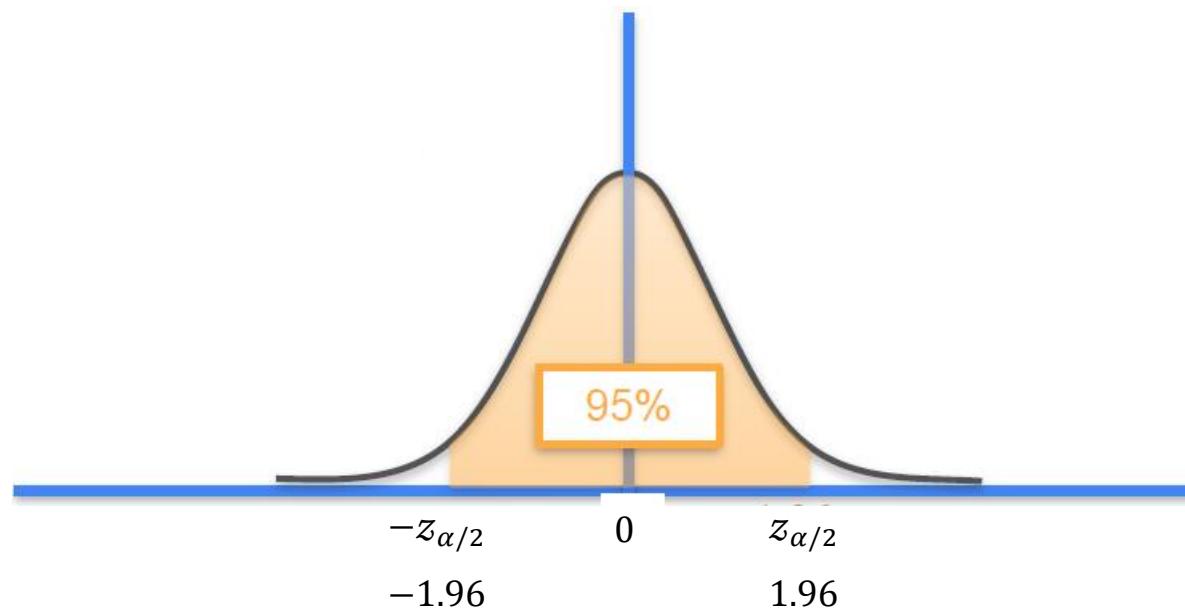


... Confidence Interval Calculation

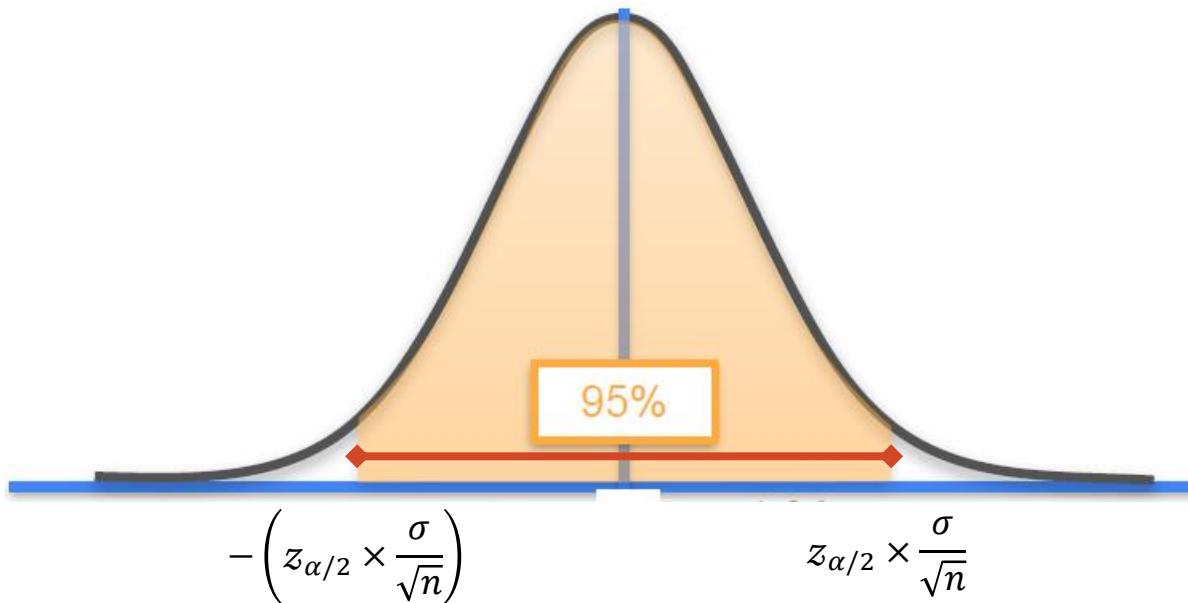
- To find the confidence interval, follow the following steps:
 1. Find the sample mean (\bar{X})
 2. Define a desired confidence level ($1 - \alpha$), e.g. 95%
 3. Get the critical value ($z_{\alpha/2}$), e.g. 1.96
 4. Find the standard error ($\frac{\sigma}{\sqrt{n}}$)
 5. Calculate the margin of error ($z_{\alpha/2} \times \frac{\sigma}{\sqrt{n}}$)
 6. Add/subtract the margin of error to/from the sample mean
 - Confidence interval = $\bar{X} \pm z_{\alpha/2} \times \frac{\sigma}{\sqrt{n}}$



... Confidence Interval Calculation



... Confidence Interval Calculation



... Confidence Interval Calculation

- Example- Average height of a city with the population 6000
 - Assumptions
 - Random sampling
 - The sample means follow the Normal distribution ($n > 30$)
 - We know that the standard deviation of the population is 25cm ($\sigma = 25$)
 - Sample size (n) is 49
 - The desired confidence level is 95%
 - Calculation
 - $\bar{X} = 170\text{cm}$
 - $z_{\alpha/2} = 1.96$
 - Margin of error = $z_{\alpha/2} \times \frac{\sigma}{\sqrt{n}} = 1.96 \times \frac{25}{7} = 7$
 - Confidence interval = $170 \pm 7 = [163, 177]$



We are 95% sure that the average height of the people in the city lies between 163 and 177cm.



Sample Size Calculation

- Example- Continued from the previous slide
 - The 7cm margin of error is a large number.
 - We want to have a smaller margin of error, say 3 cm.
 - This will give us a more exact estimate of the population mean.
 - How can we obtain a smaller confidence interval (be more confident)?
 - Use a bigger sample
 - What is the smallest sample size to obtain the desired margin of error?
 - Use the backward calculation
 - Margin of error = $z_{\alpha/2} \times \frac{\sigma}{\sqrt{n}}$
 - $3 \geq 1.96 \times \frac{25}{\sqrt{n}} \rightarrow \left(\frac{1.96 \times 25}{3}\right)^2 \leq n \rightarrow n \geq 266.78 \approx 267$



... Sample Size Calculation

- Generalized formula
 - The smallest sample size to obtain a particular margin of error, say MOE is calculated by the following formula:

$$n \geq \left(\frac{z_{\alpha/2} \times \sigma}{MOE} \right)^2$$



Unknown Standard Deviation

- Often we do not know the standard deviation of the population (σ)
 - So, we cannot calculate the Confidence interval $(\bar{X} \pm z_{\alpha/2} \times \frac{\sigma}{\sqrt{n}})$
- Solution: Use the actual sample standard deviation (s)
 - This changes the distribution shape!
 - \bar{X} no longer follows the normal distribution $\mathcal{N}\left(\mu_{\bar{X}}, \frac{\sigma^2}{n}\right)$
 - $\frac{\bar{X} - \mu}{\frac{s}{\sqrt{n}}}$ is no longer a standard Normal distribution $\mathcal{N}(0, 1^2)$, it follows a distribution called **Student's t-distribution**

$$\begin{array}{ccc} \sigma & \xrightarrow{\hspace{2cm}} & s \\ \frac{\bar{X} - \mu}{\sigma} & \xrightarrow{\hspace{2cm}} & \frac{\bar{X} - \mu}{\frac{s}{\sqrt{n}}} \end{array}$$

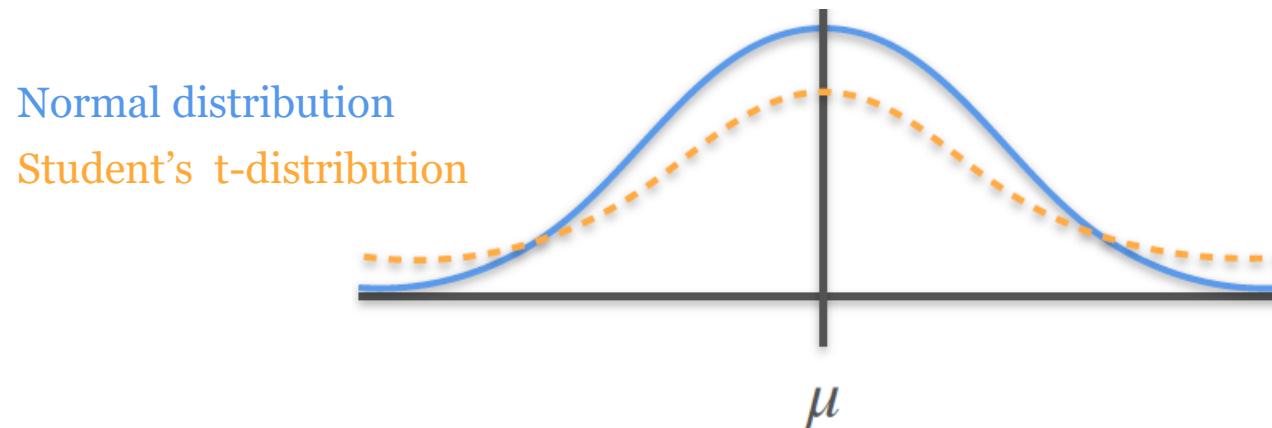
Normal distribution

Student's t-distribution



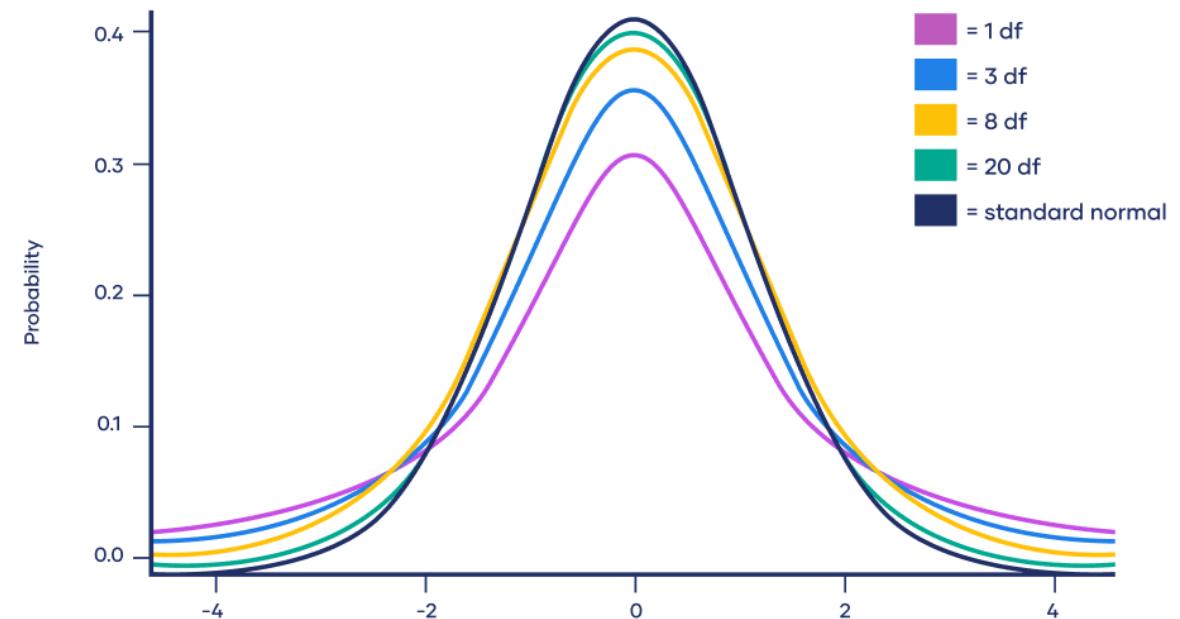
... Unknown Standard Deviation

- Student's t-distribution
 - A continuous probability distribution that generalizes the standard Normal distribution.
 - Like the standard Normal distribution, it is symmetric around zero and bell-shaped
 - In contrast to the Normal distribution, it has thicker tails
 - It means that if you sample a point out of the Student t-distribution, it is more likely to be far from the center than you pick it from the Normal distribution.

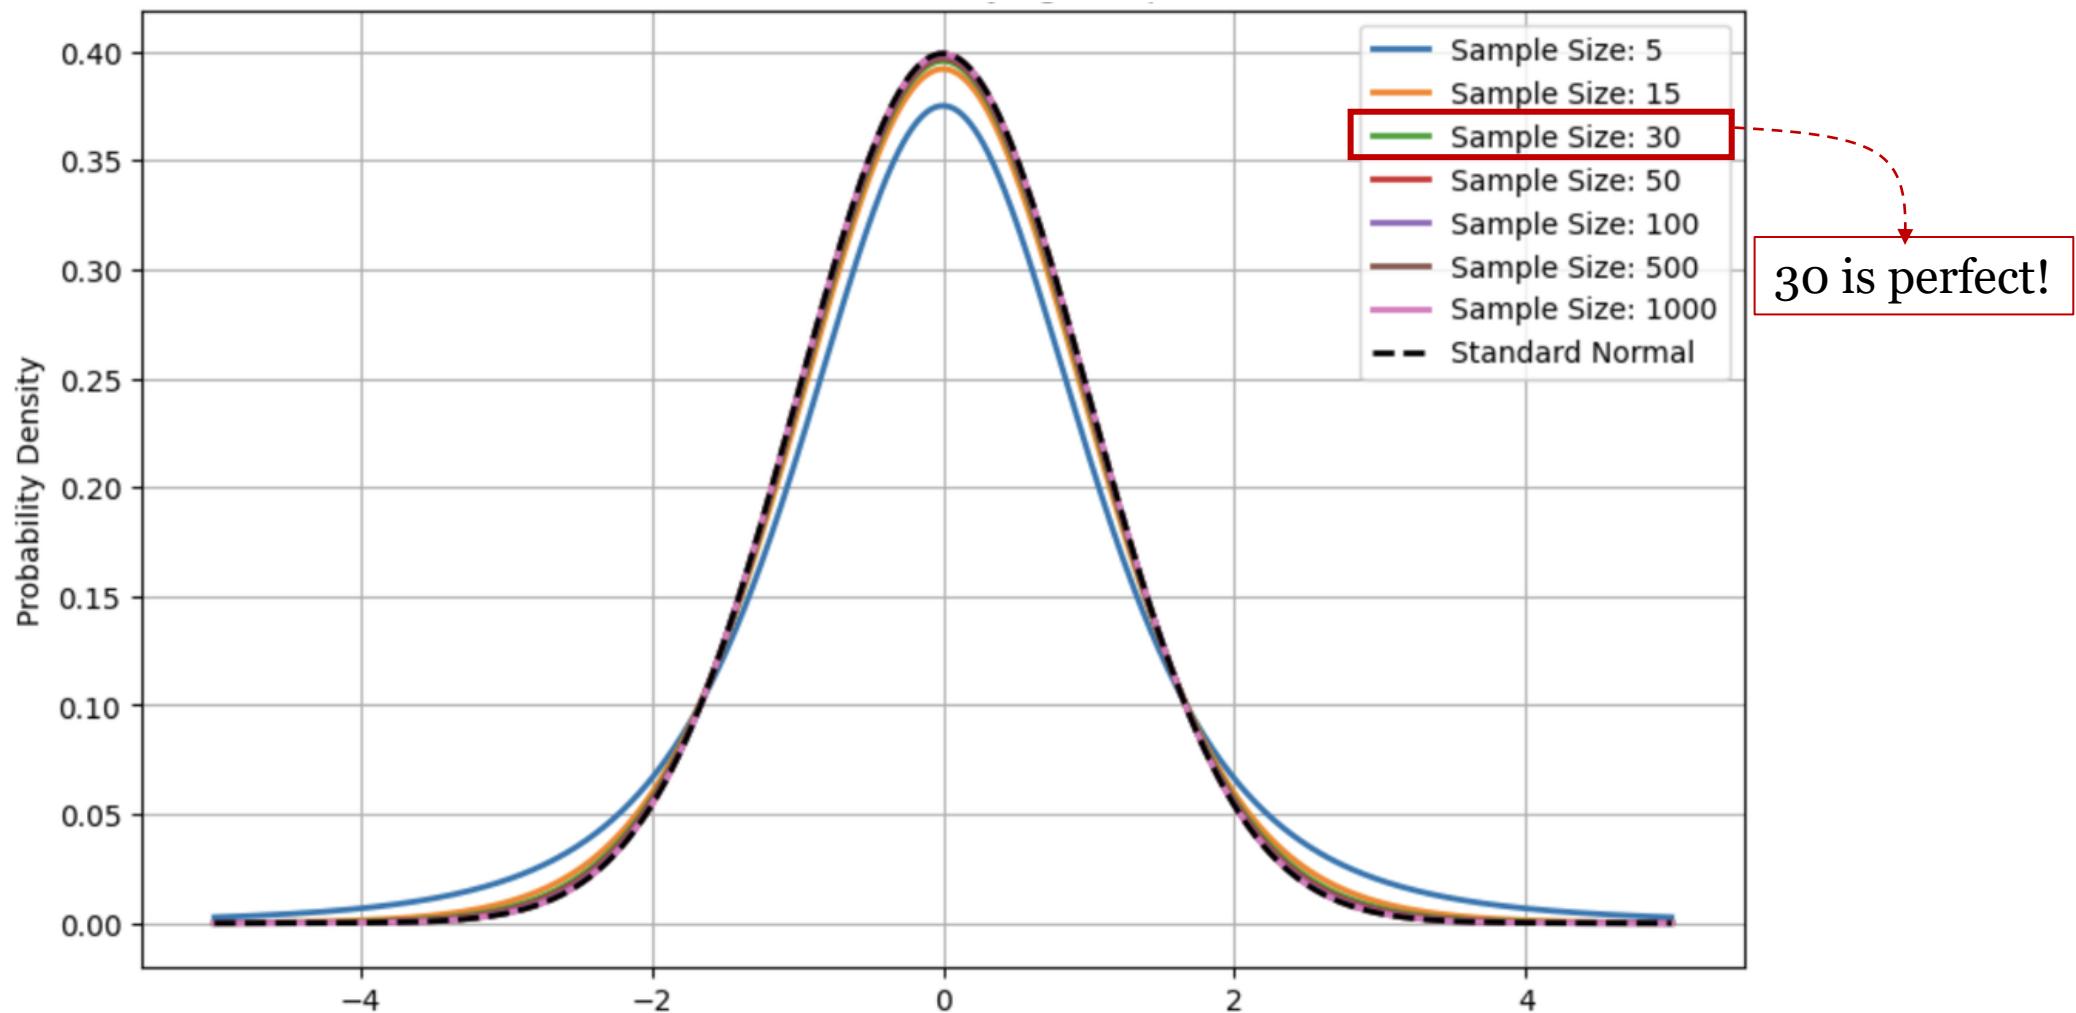


... Unknown Standard Deviation

- ... Student's t-distribution
 - It has a parameter called the degrees of freedom and shown by ν
 - $\nu = n - 1$
 - The larger the number of degrees of freedom, the closer you get to a Normal distribution
 - For $\nu = 1$, the Student's t-distribution has heavy tails
 - When $\nu \rightarrow \infty$, the Student's t-distribution becomes the standard Normal distribution $\mathcal{N}(0, 1)$



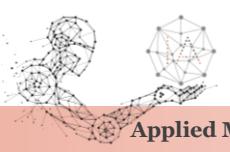
... Unknown Standard Deviation



... Unknown Standard Deviation

$1 - \alpha$

ν	50%	60%	70%	80%	90%	95%	98%	99%	99.5%	99.8%	99.9%
1	1.000	1.376	1.963	3.078	6.314	12.706	31.821	63.657	127.321	318.309	636.619
2	0.816	1.061	1.386	1.886	2.920	4.303	6.965	9.925	14.089	22.327	31.599
3	0.765	0.978	1.250	1.638	2.353	3.182	4.541	5.841	7.453	10.215	12.924
4	0.741	0.941	1.190	1.533	2.132	2.776	3.747	4.604	5.598	7.173	8.610
5	0.727	0.920	1.156	1.476	2.015	2.571	3.365	4.032	4.773	5.893	6.869
6	0.718	0.906	1.134	1.440	1.943	2.447	3.143	3.707	4.317	5.208	5.959
7	0.711	0.896	1.119	1.415	1.895	2.365	2.998	3.499	4.029	4.785	5.408
8	0.706	0.889	1.108	1.397	1.860	2.306	2.896	3.355	3.833	4.501	5.041
9	0.703	0.883	1.100	1.383	1.833	2.262	2.821	3.250	3.690	4.297	4.781
10	0.700	0.879	1.093	1.372	1.812	2.228	2.764	3.169	3.581	4.144	4.587
...
30	0.683	0.854	1.055	1.310	1.697	2.042	2.457	2.750	3.030	3.385	3.646
...
∞	0.674	0.842	1.036	1.282	1.645	1.960	2.326	2.576	2.807	3.090	3.291



... Unknown Standard Deviation

- ... Solution
 - Instead of $\bar{X} \pm z_{\alpha/2} \times \frac{\sigma}{\sqrt{n}}$ use the following formula to calculate the confidence interval:

$$\bar{X} \pm t_{\alpha/2} \times \frac{s}{\sqrt{n}}$$

t-score

How to find?

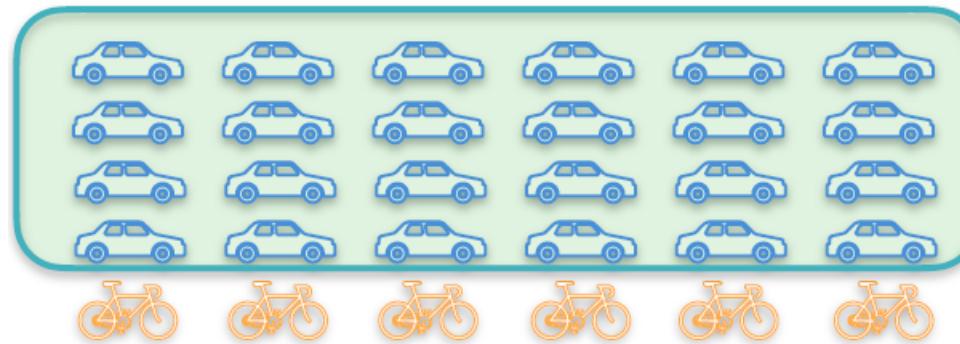
↳ Lookup in the t-table



Confidence Interval for Proportions

- Example- Proportion of people who use car

- $\hat{p} = \frac{x}{n} = \frac{24}{30} = 80\%$



- How do you calculate a 95% confidence interval for this sample proportion?

Confidence interval = $\hat{p} \pm \text{margin of error}$

$$\text{margin of error} = z_{\alpha/2} \times \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$$



... Confidence Interval for Proportions

- ... Example
 - The desired confidence level is 95%
 - $\hat{p} = 0.8$
 - $z_{\alpha/2} = 1.96$
 - Margin of error = $z_{\alpha/2} \times \sqrt{\frac{\hat{p}(1-\hat{p})}{n}} = 1.96 \times \sqrt{\frac{0.8(1-0.8)}{30}} = 0.14$
 - Confidence interval = $0.8 \pm 0.14 = [0.66, 0.94]$
 - $66\% < p < 94\%$



... Confidence Interval for Proportions

Confidence Interval for Means

Confidence interval = $\bar{X} \pm$ margin of error

$$\text{margin of error} = z_{\alpha/2} \times \frac{\sigma}{\sqrt{n}}$$

Confidence Interval for Proportions

Confidence interval = $\hat{p} \pm$ margin of error

$$\text{margin of error} = z_{\alpha/2} \times \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$$

